

NEXORA · BACKEND ENGINEERING

Nexora Backend Technical Documentation

Architecture, APIs, Business Workflows & Integrations

Priority section: ChannelEngine Compatibility for StockConnect

Version	1.0
Date	29 September 2026
Repository	nexora-backend
Runtime	Node.js 24+, Fastify 5, PostgreSQL, Redis, BullMQ
Classification	Internal engineering documentation
Audience	Backend engineers, architects, QA, DevOps, technical leads

Table of contents

1	Document control
1.1	Version history
1.2	How to read this document
2	Executive summary
2.1	The ChannelEngine relationship, stated precisely
2.2	Capability summary
2.3	What this document is for
3	Nexora backend overview
3.1	Business purpose
3.2	Deployment shape
3.3	Design principles observed in the code
4	System architecture
4.1	Layered structure
4.2	Enforced architectural rules
4.3	Module dependency map
4.4	Request path versus background path
5	Technology stack
5.1	Runtime and framework
5.2	Validation and API documentation
5.3	Data and messaging
5.4	Security
5.5	Observability
5.6	Development and quality tooling
5.7	Technologies deliberately absent
6	Repository structure
6.1	Top level
6.2	Application shell — <code>src/app/</code>
6.3	Modules — <code>src/modules/</code>
6.4	Shared utilities — <code>src/shared/</code>
7	Application lifecycle
7.1	API process startup
7.2	Worker process startup
7.3	Graceful shutdown
7.4	Request lifecycle
7.5	Plugin and route registration order
7.6	Readiness probes
8	Functional modules

8.1	Identity
8.2	Authorization
8.3	API keys
8.4	MFA
8.5	Tenants
8.6	Audit
8.7	Products
8.8	Offers
8.9	Pricing
8.10	Inventory
8.11	Orders
8.12	Shipments
8.13	Returns
8.14	Cancellations
8.15	Channels and marketplaces
8.16	Webhooks
8.17	External ID mapping
9	API architecture
9.1	Versioning and surfaces
9.2	Validation and serialisation
9.3	Success and error envelopes
9.4	Error taxonomy
9.5	Idempotency
9.6	Pagination, filtering and rate limiting
9.7	API documentation endpoints
10	API reference
10.1	Surface totals
10.2	Operational and documentation endpoints
10.3	Identity, tenancy and access control
10.4	Catalogue, pricing and inventory
10.5	Orders, shipments, returns and cancellations
10.6	Channels, marketplaces and webhooks
10.7	Compatibility surfaces
10.8	Worker observability endpoints
11	Authentication and authorization
11.1	Credential types
11.2	The public route allowlist
11.3	Authorization flow
11.4	Secret storage
11.5	Permission catalogue

11.6 Verified security limitations

12 Database architecture

12.1 Technology and access pattern

12.2 Migrations

12.3 Core entity relationships

12.4 Identity and access tables

12.5 Table inventory

12.6 Tenant isolation mechanics

12.7 Data retention

13 ChannelEngine integration for StockConnect

13.1 What the integration actually is

13.2 System context

13.3 Responsibility split

13.4 Compatibility module architecture

13.5 Authentication for compatibility clients

13.6 Endpoint inventory

13.7 Identifier mapping

13.8 Idempotency

13.9 Error mapping

13.10 Product data flow

13.11 Inventory data flow

13.12 Pricing data flow

13.13 Order lifecycle

13.14 Delivery-state reporting

13.15 Invoice retrieval

13.16 Outbound webhooks to StockConnect

13.17 Readiness for the compatibility surface

13.18 Operational workflows

13.19 Troubleshooting

13.20 Implementation status

13.21 Future considerations

14 Product, inventory, pricing and order workflows

14.1 Publishing a product to a marketplace

14.2 Inventory movement through an order's life

14.3 Channel-fulfilled orders

14.4 Request equivalence in practice

15 Queue and background processing

15.1 Queues and jobs

15.2 Worker architecture

15.3 Retry semantics

15.4	Schedulers
16	Marketplace integrations
16.1	Adapter ports
16.2	Support matrix
16.3	Authentication per provider
16.4	Catalog synchronisation pipeline
16.5	Order ingestion
16.6	Inbound webhook flow
17	Webhooks
17.1	Subscription model
17.2	Event catalogue
17.3	Delivery
17.4	SSRF protection
17.5	Secret rotation
18	External services
18.1	Outbound HTTP client
18.2	A configuration inconsistency worth knowing
19	Error handling and observability
19.1	Error normalisation
19.2	Structured logging
19.3	Metrics
19.4	Tracing
19.5	Health and readiness
19.6	Diagnostic workflow
20	Development environment
20.1	Prerequisites
20.2	First run
20.3	Required configuration
20.4	Configuration reference by concern
20.5	npm scripts
20.6	Common local problems
21	Testing strategy
21.1	Suites
21.2	Coverage by area
21.3	Running tests
21.4	Verified results
21.5	Gaps
22	Deployment and operations
22.1	Container image
22.2	Process topology

22.3	Probe configuration
22.4	Migrations in production
22.5	Shutdown behaviour
22.6	Security hardening checklist for production
23	Troubleshooting
23.1	Authentication and authorization
23.2	Inventory and orders
23.3	Synchronisation
23.4	Webhooks
23.5	Platform
24	Known limitations
24.1	Correctness and completeness
24.2	Compatibility surface
24.3	Security
24.4	Platform and process
25	Future technical considerations
26	Glossary
27	Appendix A — Feature to source map
28	Appendix B — Architecture decision records
29	Appendix C — Related documentation

1 Document control

Field	Value
Document title	Nexora Backend Technical Documentation
Subtitle	Architecture, APIs, Business Workflows & Integrations
Focus area	ChannelEngine compatibility for StockConnect
Version	1.0
Status	Issued
Classification	Internal engineering documentation
Source repository	nexora-backend
Analysed branch	docs/api-audit-swagger-postman
Analysed commit base	5e610ec (dev) plus documentation branch
Runtime analysed	Node.js 24+, Fastify 5, PostgreSQL, Redis, BullMQ
Intended audience	Backend engineers, architects, QA, DevOps, technical leads, integration partners

1.1 Version history

Version	Date	Author	Change
1.0	2026-09-29	Engineering documentation review	First complete technical reference produced from a full source-code audit of the repository.

1.2 How to read this document

Statements in this document are grouped into three categories, and the wording makes the category explicit:

Category	Wording used	Meaning
Verified from source	"Nexora does X", with a file reference	Behaviour was read in the implementation
Verified from project documentation	"The ADR records...", "the runbook states..."	Taken from in-repository documents, not re-derived from code
Recommendation	"Consider...", "A future change could..."	Not implemented; an opinion for planning

Where behaviour could not be verified, the document says so instead of guessing. Source references use paths relative to the repository root.

2 Executive summary

Nexora is a multi-tenant commerce and channel-management backend. It owns product, offer, price, inventory, order, shipment, return and cancellation data for each tenant, connects that data to external marketplaces, and exposes the whole surface through a versioned HTTP API.

The system is a **modular monolith** deployed as **two processes** that share one codebase and one database: an **API process** (`src/app/main.js`) that serves HTTP traffic, and a **worker process** (`src/workers/main.js`) that consumes background jobs. PostgreSQL is the single source of truth. Redis provides caching, distributed locking, rate limiting and the BullMQ queue backend. Neither Redis nor the queue holds authoritative business state.

The API exposes **147 registered endpoints** on the main server across three surfaces:

Surface	Prefix	Purpose
Native API	<code>/api/v1/*</code>	Nexora's own contract — the richest and most complete surface
Merchant compatibility	<code>/api/v2/*</code>	A ChannelEngine-shaped contract for merchant integrations
StockConnect CE compatibility	<code>/api/v2/ce/*</code>	An additive ChannelEngine-shaped contract targeted at StockConnect

The worker process exposes three further endpoints on a separate port for liveness, readiness and metrics.

2.1 The ChannelEngine relationship, stated precisely

This is the single most misunderstood point about the system, so it is stated plainly here and expanded in the dedicated chapter.

Nexora does not call ChannelEngine. Nexora replaces it.

A full-text search of `src/` finds the string "ChannelEngine" in exactly three files, all of them comments and route descriptions inside the compatibility module. There is no ChannelEngine API client, no ChannelEngine credentials in configuration, and no outbound HTTP integration to ChannelEngine anywhere in the codebase. The compatibility module's `infrastructure/` directory is an empty boundary stub.

What Nexora actually does is implement a **ChannelEngine-compatible inbound API** so that StockConnect — an existing external merchant system that was written against ChannelEngine — can be repointed at Nexora with minimal client-side change. StockConnect is the HTTP client; Nexora is the HTTP server. Data flows from StockConnect into Nexora, and from Nexora back to StockConnect through polling endpoints and outbound webhooks.

Everything downstream of that boundary — reaching the actual sales channels — is handled by Nexora's own marketplace adapters for Amazon, Noon, Namshi and Shopify, which are independent of ChannelEngine entirely.

2.2 Capability summary

Area	State
Identity, RBAC, API keys, MFA step-up	Implemented and covered by tests
Multi-tenancy with PostgreSQL row-level security	Implemented; see the noted gaps in the security chapter
Product, offer, price, inventory, order, shipment, return, cancellation	Implemented with a transactional domain model
Merchant <code>/api/v2</code> compatibility surface	Implemented, idempotency-enforced
StockConnect <code>/api/v2/ce</code> compatibility surface	Implemented; several documented partial areas
Outbound catalog, price and inventory sync to marketplaces	Implemented via queue-driven adapters
Inbound marketplace order ingestion	Implemented for Shopify by polling; lifecycle/webhook-driven for Amazon, Noon, Namshi
Outbound tenant webhooks with HMAC signing and retry	Implemented
Observability: structured logs, Prometheus metrics, OpenTelemetry tracing	Implemented
Live StockConnect staging or production verification	Not performed in-repository ; the cutover runbook states this explicitly

2.3 What this document is for

An engineer who reads it should be able to locate any subsystem in the source tree, follow a request from the socket to the database and back, understand how a marketplace order becomes a Nexora order, diagnose a failed synchronisation, and continue development without needing the original authors.

3 Nexora backend overview

3.1 Business purpose

A tenant using Nexora is a merchant or merchant group that sells the same catalogue across several marketplaces. The backend solves four problems for them.


One catalogue, many channels. Products are defined once per tenant with a stable merchant SKU. Each combination of product and sales channel becomes an *offer*, which carries the channel-specific listing state and external identifier. Prices are separate records with currency, validity windows and optional channel targeting. This separation means a price change or a channel-specific listing change does not mutate the product record.

Accurate stock. Inventory is authoritative in PostgreSQL, held per stock location, and split into `on_hand`, `reserved` and `available` with the invariant `available = on_hand - reserved` enforced in the domain layer. Orders reserve stock at creation; shipments consume it; returns restore it. Every change writes an append-only movement row, so stock history is auditable.

Order lifecycle across channels. Orders arrive natively through the API, through channel ingestion endpoints, by polling a marketplace, or through a marketplace webhook. All paths converge on the same domain model, the same status machine, and the same inventory effects. Shipments, returns and cancellations hang off the order and update line-level quantities.

Integration without coupling. The core domain modules know nothing about ChannelEngine, Amazon or Shopify. Provider-specific code lives behind adapter ports, and the ChannelEngine-shaped API lives in a separate `compatibility` module. This boundary is not merely a convention — it is mechanically enforced by dependency-cruiser rules that fail the build when a core module imports the compatibility module or when the compatibility module reaches past another module's public contract.

3.2 Deployment shape



The two processes are deliberately separate. The API process must stay responsive, so anything slow, retryable, or dependent on a third party is pushed onto a queue and executed by the worker. The API process starts the outbox publisher; the worker process starts the retention cleanup and catalog reconciliation schedulers. Both connect to the same PostgreSQL and Redis instances.


3.3 Design principles observed in the code

Principle	How it is realised
PostgreSQL is authoritative	ADR-002; all durable state in Postgres, Redis never holds business truth
Layered modules	Every module under <code>src/modules/</code> has <code>domain/</code> , <code>application/</code> , <code>infrastructure/</code> , <code>presentation/</code> , <code>public/</code>
Cross-module access through contracts	Modules import each other only via <code>public/</code> ; enforced by <code>no-cross-module-internals</code>
Provider neutrality in the core	<code>core-no-compatibility</code> and <code>compatibility-public-contracts-only</code> rules
Domain purity	<code>domain-no-infrastructure-packages</code> forbids Fastify, pg, ioredis, bullmq, zod inside <code>domain/</code>
Transactional consistency	Domain mutation and outbox event write share one transaction
Explicit composition	No DI container; dependencies are constructed in <code>create-infrastructure.js</code> and <code>create-application.js</code> and passed down

4 System architecture

4.1 Layered structure

Nexora follows a hexagonal, layered arrangement inside each module. The direction of dependency always points inward toward the domain.



The `domain/` layer holds business rules with no framework imports at all. The `application/` layer orchestrates: it loads aggregates through repository ports, applies domain operations, records integration events, and commits. The `infrastructure/` layer implements the ports with raw SQL through `pg`. The `presentation/` layer is thin — it validates with `Zod`, calls one use case, and shapes the response. The `public/` layer is the only surface other modules are allowed to import.


4.2 Enforced architectural rules

These are not guidelines; `npm run arch:check` fails the build on violation. They are defined in `.dependency-cruiser.cjs`.

Rule	Severity	What it prevents
<code>no-circular</code>	error	Circular dependencies anywhere in <code>src/</code>
<code>domain-no-infrastructure-layer</code>	error	Domain importing <code>src/infrastructure</code> or <code>src/app</code>
<code>domain-no-infrastructure-packages</code>	error	Domain importing Fastify, pg, ioredis, bullmq, AWS SDK, pino, prom-client, OpenTelemetry or zod
<code>domain-no-module-infrastructure</code>	error	Domain importing its own module's infrastructure or presentation
<code>application-no-presentation</code>	error	Application depending on presentation
<code>no-cross-module-internals</code>	error	A module reaching into another module's <code>domain/</code> , <code>infrastructure/</code> or <code>presentation/</code>
<code>core-no-compatibility</code>	error	Any non-compatibility module importing the compatibility module
<code>compatibility-no-direct-persistence</code>	error	The compatibility module using Postgres repositories directly
<code>compatibility-public-contracts-only</code>	error	The compatibility module bypassing other modules' <code>public/</code> contracts
<code>shared-stays-generic</code>	error	<code>src/shared/</code> depending on app, infrastructure, modules or workers
<code>infrastructure-no-modules</code>	error	<code>src/infrastructure/</code> depending on modules or app
<code>no-dev-dep-in-src</code>	error	Runtime code importing a devDependency
<code>no-orphans</code>	warn	Unreachable files, excluding the two process entrypoints

The `core-no-compatibility` rule is the structural guarantee behind the claim that Nexora's domain is provider-neutral. The ChannelEngine-shaped API can be deleted without touching a single core module.

4.3 Module dependency map




Arrows show the direction of the import, always through a `public/` contract. Note that no arrow points *into* `compatibility` — that is the enforced rule made visible.

4.4 Request path versus background path

The system has two distinct execution paths, and understanding which one a piece of work is on explains most of its behaviour.

Synchronous HTTP path. A client request is validated, authenticated, authorised, and executed inside a database transaction. The transaction writes the domain change and, in the same commit, appends an integration event row to the outbox. The response returns immediately after commit. Nothing that requires a third-party HTTP call happens here.

Asynchronous background path. The outbox publisher polls for unpublished events and enqueues BullMQ jobs. The worker consumes them, and only at that point does outbound HTTP occur — a catalog sync call to a marketplace, a webhook delivery to a tenant endpoint, or a marketplace lifecycle operation. Failures here are retried with backoff and never affect the original API response.



This is the transactional outbox pattern from ADR-005. Its value is that an event is never lost when the process crashes between the database commit and the queue publish, because the event is part of the commit.

5 Technology stack

Only technologies actually present in `package.json` and used in `src/` are listed. Purpose and location are given for each.

5.1 Runtime and framework

Technology	Version constraint	Role in Nexora	Where
Node.js	<code>>=24.0.0</code> (engines)	Runtime for both processes	<code>package.json</code>
Fastify	<code>^5.2.0</code>	HTTP server, plugin system, lifecycle hooks	<code>src/app/http/create-server.js</code>
<code>fastify-plugin</code>	<code>^5.0.1</code>	Registers plugins on the root instance so hooks apply globally	Auth, logging, metrics plugins
JavaScript (ESM)	<code>"type": "module"</code>	Entire codebase; no TypeScript build step	Throughout

Nexora is plain ESM JavaScript. There is no transpilation: `npm run build` runs `node --check` over every file as a syntax gate, nothing more.

5.2 Validation and API documentation

Technology	Role	Where
Zod <code>^3.24.1</code>	Every request and response schema; also configuration validation	<code>presentation/*.schemas.js</code> , <code>src/app/config/schema.js</code>
<code>fastify-type-provider-zod</code> <code>^4.0.2</code>	Binds Zod schemas to Fastify's validator and serializer compilers	<code>create-server.js</code>
<code>@fastify/swagger</code> <code>^9.4.0</code>	Generates the OpenAPI 3.1 document from those same Zod schemas	<code>create-server.js</code>
<code>@scalar/fastify-api-reference</code> <code>^1.25.116</code>	Serves the interactive API reference UI at <code>/docs</code>	<code>create-server.js</code>

This chain matters: because the OpenAPI document is generated from the Zod schemas that Fastify actually enforces at runtime, the documentation cannot drift from validation. A schema change updates both simultaneously.

5.3 Data and messaging

Technology	Role	Where
PostgreSQL via <code>pg</code> <code>^8.13.1</code>	Source of truth; raw SQL in repositories, no ORM (ADR-011)	<code>src/infrastructure/postgres/</code>
Redis via <code>ioredis</code> <code>^5.4.1</code>	Cache, distributed lock, GCRA rate limiter, BullMQ transport	<code>src/infrastructure/redis/</code>
BullMQ <code>^5.34.0</code>	Four job queues with retry and backoff	<code>src/infrastructure/queue/</code>
<code>@aws-sdk/client-s3</code> <code>^3.700.0</code>	S3-compatible object storage provider	<code>src/infrastructure/storage/</code>

Object storage is wired at bootstrap and health-checked by the API readiness probe, but no domain module currently writes objects through it. It is infrastructure-ready rather than in use.

5.4 Security

Technology	Role	Where
<code>argon2</code> <code>^0.45.1</code>	Argon2id password hashing	<code>src/infrastructure/auth/argon2-password-hasher.js</code>
<code>jose</code> <code>^6.2.12</code>	JWT signing and verification, HS256 or RS256	<code>src/infrastructure/auth/jose-access-token-service.js</code>
<code>otplib</code> <code>^13.5.0</code>	TOTP enrolment and verification for MFA step-up	<code>src/modules/mfa/</code>
<code>@fastify/helmet</code> <code>^12.0.1</code>	Security response headers; CSP explicitly disabled	<code>create-server.js</code>
<code>@fastify/cors</code> <code>^10.0.1</code>	CORS, enabled only when an origin allowlist is configured	<code>create-server.js</code>
Node <code>crypto</code>	SHA-256 token hashing, AES-256-GCM secret encryption, HMAC webhook signing	<code>src/shared/security/</code>

5.5 Observability

Technology	Role	Where
<code>pino</code> <code>^9.5.0</code>	Structured JSON logging with redaction and request-context enrichment	<code>src/infrastructure/observability/pino-logger.js</code>
<code>prom-client</code> <code>^15.1.3</code>	Prometheus metrics, including Node default metrics	<code>src/infrastructure/observability/prometheus-metrics.js</code>
OpenTelemetry SDK <code>^0.57.0</code>	Distributed tracing with HTTP, <code>pg</code> and <code>ioredis</code> instrumentation	<code>src/infrastructure/observability/tracing.js</code>

5.6 Development and quality tooling

Technology	Role
Vitest <code>^3.2.7</code>	Unit and integration test runner, two named projects
ESLint <code>^9.17.0</code>	Linting, including a rule restricting direct <code>process</code> access outside configuration and endpoints
<code>dependency-cruiser</code> <code>^16.8.0</code>	Enforces the architectural boundary rules listed earlier
Prettier <code>^3.4.2</code>	Formatting; not part of the <code>verify</code> gate
<code>dotenv</code> <code>^18.0.1</code>	Loads <code>.env</code> in integration tests and tooling scripts only

5.7 Technologies deliberately absent

Documenting absences prevents wrong assumptions.

Not present	Consequence
No ORM	All persistence is hand-written SQL in repository classes
No TypeScript	Type safety comes from Zod at boundaries, not from compilation
No ChannelEngine SDK or client	Nexora never calls ChannelEngine; see the compatibility chapter
No GraphQL	The API is REST-style JSON over HTTP
No CI workflow in the repository	<code>.github/workflows/</code> does not exist; quality gates are run manually or by external tooling

6 Repository structure

6.1 Top level

Path	Contents
<code>src/app/</code>	Process bootstrap, configuration, HTTP server, plugins, operational routes, readiness
<code>src/modules/</code>	Twenty business modules, each layered
<code>src/infrastructure/</code>	Technology adapters: Postgres, Redis, queue, storage, HTTP client, observability, auth primitives
<code>src/shared/</code>	Cross-cutting utilities with no dependency on modules or infrastructure
<code>src/workers/</code>	Worker process entryptpoint, job handlers, worker bootstrap wiring, worker observability server
<code>tests/unit/</code>	136 test files, no external dependencies
<code>tests/integration/</code>	73 test files, requiring PostgreSQL, Redis and MinIO
<code>docs/</code>	Architecture notes, ADRs, operations runbooks, API inventories, this document
<code>scripts/</code>	Documentation and maintenance tooling
<code>infrastructure/</code>	Local observability stack configuration for Docker Compose

The source tree contains 826 JavaScript files under `src/`.

6.2 Application shell — `src/app/`

Path	Responsibility
<code>main.js</code>	API process entryptpoint: config, infrastructure, application, listen, signal handlers
<code>bootstrap/create-infrastructure.js</code>	Constructs logger, metrics, tracing, Postgres pool, migrations, Redis, queue, storage, outbox, idempotency
<code>bootstrap/create-application.js</code>	Constructs every module, wires use cases, registers readiness probes, builds the HTTP server
<code>bootstrap/shutdown.js</code>	Ordered graceful shutdown with a global deadline
<code>bootstrap/migrate-cli.js</code>	Standalone <code>up</code> and <code>status</code> migration commands
<code>bootstrap/external-id-backfill-cli.js</code>	Backfills compatibility integer identifiers for existing rows
<code>config/schema.js</code>	Zod schema for every recognised environment variable
<code>config/config.js</code>	Maps validated environment into the nested <code>AppConfig</code> and applies cross-field consistency checks

Path	Responsibility
<code>http/create-server.js</code>	Fastify construction and the full plugin and route registration order
<code>http/plugins/</code>	<code>request-context</code> , <code>logging</code> , <code>metrics</code> , <code>error-handler</code> , <code>authentication</code>
<code>http/routes/</code>	<code>health.routes.js</code> , <code>metrics.routes.js</code> , <code>foundation.routes.js</code>
<code>observability/readiness.js</code>	The readiness service and the default and worker probe sets
<code>errors/</code>	<code>error-mapper.js</code> and <code>error-envelope.js</code> — normalisation and response shaping

6.3 Modules — `src/modules/`

Module	Concern
<code>identity</code>	Login, refresh rotation, logout, access-token authentication
<code>authorization</code>	Roles, permissions, memberships, effective permission resolution
<code>api-keys</code>	API key issuance, rotation, revocation, verification
<code>mfa</code>	TOTP enrolment, step-up verification, recovery codes
<code>tenants</code>	Tenant lifecycle and slug rules
<code>audit</code>	Append-only audit log and its query API
<code>products</code>	Product catalogue and localised content
<code>offers</code>	Product-to-channel listings
<code>pricing</code>	Prices with currency and validity windows
<code>inventory</code>	Stock locations, balances, reservations, movements
<code>orders</code>	Order capture, status machine, fulfilment coordination
<code>shipments</code>	Shipment documents and inventory consumption
<code>returns</code>	Return requests, approval, receipt, restock
<code>cancellations</code>	Line cancellation and reservation release
<code>channels</code>	Sales channels and their marketplace connections
<code>marketplaces</code>	Marketplace registry and all provider adapters
<code>channel-catalog-sync</code>	Outbound product, offer, price and inventory synchronisation
<code>marketplace-order-ingestion</code>	Inbound order ingestion and lifecycle processing
<code>marketplace-webhook-ingestion</code>	Inbound marketplace webhook receipt and dispatch

Module	Concern
webhooks	Outbound tenant webhook subscriptions and delivery
external-id-mapping	Stable integer identifiers for compatibility consumers
compatibility	The ChannelEngine-shaped <code>/api/v2</code> and <code>/api/v2/ce</code> surfaces


6.4 Shared utilities — `src/shared/`

Path	Purpose
<code>context/request-context.js</code>	AsyncLocalStorage request scope carrying request id, tenant, principal, trace id
<code>idempotency/</code>	Idempotency key requirement, stable SHA-256 fingerprinting, service port
<code>errors/</code>	<code>AppError</code> hierarchy, error code registry, log-safe error inspection
<code>events/</code>	Integration event catalogue, event recorder port, external delivery allowlist
<code>money/</code>	ISO 4217 currency parsing and integer minor-unit amounts
<code>pagination/</code>	Cursor encoding and clamped page limits
<code>rate-limit/</code> , <code>auth/rate-limit-policies.js</code>	Rate limiter port and the named policies
<code>auth/merchant-compat-route-prefix.js</code>	Identifies <code>/api/v2/ce</code> paths for the CE credential reader
<code>auth/read-merchant-compat-query-api-key.js</code>	Reads <code>apiKey</code> , <code>apikey</code> query values and the <code>X-CE-KEY</code> header
<code>security/</code>	SSRF validator, outbound HTTPS validation, password policy, crypto helpers
<code>logging/redaction.js</code>	Pino redaction paths and deep redaction for arbitrary payloads
<code>stockconnect/</code>	StockConnect CE webhook subscription constants

7 Application lifecycle

7.1 API process startup

The sequence in `src/app/main.js` is strictly ordered; each step depends on the previous one.



Two details deserve emphasis. First, **migrations run automatically at infrastructure boot**. If `DATABASE_MIGRATION_URL` is set and differs from `DATABASE_URL`, migrations execute against the migration URL on a separate pool while the runtime pool uses the ordinary connection string — this is what allows the runtime role to be a restricted, RLS-bound role. Second, the **outbox publisher runs only on the API process**; the worker does not start it.

7.2 Worker process startup

`src/workers/main.js` shares steps one and two with the API but diverges after that. It passes `processKind: 'worker'` into `createInfrastructure`, which changes the OpenTelemetry service name to `WORKER_OTEL_SERVICE_NAME` or `{OTEL_SERVICE_NAME}-worker`. It then registers job handlers on the worker runtime, builds a worker-specific readiness service, optionally starts the observability HTTP server, and starts the retention cleanup and catalog reconciliation schedulers.

Concern	API process	Worker process
HTTP server on <code>SERVER_PORT</code>	Yes	No
Observability HTTP on <code>WORKER_OBSERVABILITY_PORT</code>	No	Yes, when enabled
BullMQ workers consuming jobs	No	Yes
Outbox publisher	Yes	No
Retention cleanup scheduler	No	Yes
Catalog sync reconciliation scheduler	No	Yes
Readiness probe set	postgres, redis, queue, storage, CE compatibility	postgres, redis, queue, workers_registered

The storage probe exists only on the API side; the worker instead asserts that job handlers are registered.

7.3 Graceful shutdown

`src/app/bootstrap/shutdown.js` implements a single ordered routine used by both processes, bounded by `SERVER_SHUTDOWN_TIMEOUT_MS` (default 15000). Each step is wrapped so a failure is logged as a warning and does not abort the remaining steps, and each is raced against the time remaining until the global deadline.


1. Log shutdown start.
2. `readiness.markNotReady()` — readiness immediately reports failure so load balancers stop routing traffic.
3. Close the API HTTP server, if present.
4. Close the worker observability HTTP server, if present.
5. Stop the outbox publisher (API).
6. Stop the retention cleanup scheduler (worker).
7. Stop the catalog sync reconciliation scheduler (worker).
8. Close the worker runtime, draining in-flight jobs (worker).

9. Close the queue.
10. Close Redis.
11. Close the database pool.
12. Shut down tracing.
13. Flush the logger.

The order is the reverse of acquisition, with one deliberate exception: readiness is marked unhealthy *first*, before anything is torn down, so that in-flight requests can complete while new ones are routed elsewhere.

One asymmetry is worth noting. The worker's `/health/live` returns 503 with `{ status: 'shutting_down' }` once shutdown begins, whereas the API's `/health/live` always returns `{ status: 'ok' }` and is unaware of the shutdown flag. Orchestrators relying on API liveness to detect draining will not see it there; they should use `/health/ready`, which does reflect it.

7.4 Request lifecycle



Because the authentication plugin is registered with `fastify-plugin`, its `preHandler` hook is attached to the **root instance** and therefore applies to every route, including routes registered earlier in the file. Registration order controls which routes exist, not which are protected; protection is controlled by the `public-route` allowlist inside the plugin.

7.5 Plugin and route registration order

Order	Component	Mechanism	Purpose
1	<code>@fastify/helmet</code>	middleware	Security headers, CSP disabled
2	<code>@fastify/cors</code>	middleware	Only when <code>SERVER_CORS_ORIGINS</code> is non-empty
3	<code>request-context</code>	<code>onRequest</code>	Request id, AsyncLocalStorage, trace id
4	<code>logging</code>	<code>onResponse</code>	Structured access log
5	<code>metrics</code>	<code>onResponse</code>	HTTP counter and histogram, when metrics enabled
6	<code>error-handler</code>	<code>setErrorHandler</code>	Error normalisation and envelope
7	<code>@fastify/swagger</code>	spec generation	OpenAPI 3.1 from Zod
8	Scalar and doc routes	routes	<code>/docs</code> , <code>/openapi.json</code> , <code>/api-docs</code> , <code>/api-docs.json</code> when <code>DOCS_ENABLED</code>
9	health routes	routes	<code>/health/live</code> , <code>/health/ready</code>
10	metrics routes	routes	<code>/internal/metrics</code>
11	foundation routes	routes	<code>/api/v1/foundation/*</code>
12	tenants routes	routes	Tenant bootstrap, partly public
13	identity auth routes	routes	Login, refresh, logout
14	marketplace webhook ingestion	routes	Inbound marketplace webhooks
15	authentication	<code>preHandler</code> (global)	JWT, API key, CE key resolution
16+	authorization, audit, api-keys, mfa, marketplaces, channels, products, pricing, offers, inventory, orders, shipments, returns, cancellations, webhooks, compatibility	routes	The protected surface

7.6 Readiness probes

Probe	Checks	API	Worker
<code>postgres</code>	<code>database.healthCheck()</code>	Yes	Yes
<code>redis</code>	<code>redis.healthCheck()</code>	Yes	Yes

Probe	Checks	API	Worker
queue	queue.healthCheck()	Yes	Yes
storage	storage.healthCheck()	Yes	No
workers_registered	Job handlers registered on the runtime	No	Yes
external_compat_ce_routes	CE handler dependencies wired	Yes	No
stockconnect_ce_compat	Same check under the legacy name	Yes	No

The two compatibility probes are **synchronous wiring assertions**, not connectivity checks. They confirm that `stockConnectCeOrderCompatibilityQuery`, `stockConnectCeCatalogCommand`, `stockConnectCeOrderInvoiceQuery` and `stockConnectCeChannelCompatibilityQuery` are present on the route dependencies. They make no outbound HTTP call and say nothing about StockConnect's own availability.

When `markNotReady()` has been called, `evaluate()` short-circuits to `{ ready: false, checks: { process: 'failed' } }` without running any dependency probe.

8 Functional modules

Each module below is described by what it is for, what it owns, and where its limits are. Endpoint lists are complete for the module; full request and response schemas live in the generated OpenAPI document and are not duplicated here.

8.1 Identity

Purpose. Authenticate a user into a tenant and issue short-lived access tokens with rotating refresh tokens.

Login (`src/modules/identity/application/use-cases/login.js`) resolves the tenant by slug, requires tenant status `ACTIVE`, normalises the email through the `Email` value object, verifies the password with Argon2id, and requires an active membership that grants access. On success it creates a refresh session, signs a JWT, and writes a `LOGIN_SUCCESS` audit entry. Failures write `LOGIN_FAILURE` and return a generic `InvalidCredentialsError` so the caller cannot distinguish a wrong password from a missing user.

Refresh rotation is the most security-sensitive part of the module. The refresh token is hashed with SHA-256 before lookup, so the raw token is never stored. `evaluateRefreshSession` in `domain/refresh-rotation.js` classifies the presented token. If a session that was already `REPLACED` or `REVOKED` is presented again, this is treated as **token family reuse**: the entire `familyId` is revoked, a `REFRESH_REUSE_DETECTED` audit entry is written, and the caller receives invalid credentials. A normal rotation creates a new session in the same family and marks the old one replaced.

Access-token authentication (`application/authenticate-access-token.js`) verifies the JWT, loads the refresh session and requires it to be `ACTIVE` and unexpired, loads the user, and resolves effective permissions through the authorization module. A user with **zero** effective permissions is rejected with an authentication error rather than being admitted with an empty permission set.

Endpoints. `POST /api/v1/auth/login`, `POST /api/v1/auth/refresh`, `POST /api/v1/auth/logout`, `GET /api/v1/auth/me`. The first three are public; `me` requires a bearer token.

Limitation. Login does not enforce MFA even when a user has an active TOTP factor. MFA in Nexora gates *step-up* operations, not initial sign-in.

8.2 Authorization

Purpose. Decide what an authenticated principal may do inside a tenant.

The model is: a global `permissions` catalogue, tenant-scoped `roles`, a `role_permissions` join, and `membership_roles` linking a tenant membership to roles. Effective permissions (`domain/effective-permissions.js`) are the sorted union of permission keys from **active** roles on an **active** membership.

There is an important asymmetry between seeding and checking. Seeding system roles uses **pattern matching** — `all`, `exact`, `prefix` and `suffix` patterns in `domain/permission-pattern.js` — to expand a role template into concrete permission keys. Runtime checks in `application/authorization-service.js` are **exact string comparison** only. Wildcards never apply at request time.

There is no global authorization Fastify plugin. Each route handler obtains the actor context and calls `requirePermission` with a literal permission string. This makes authorization explicit and greppable at the cost of being easy to omit in a new handler.

The **last-admin guard** (`application/last-admin-guard.js`) prevents removing a role granting `tenant.admin` when that would leave the tenant with no admin-capable membership, using `SELECT ... FOR UPDATE` to make the check race-free.

Endpoints. `GET/POST /api/v1/roles`, `GET /api/v1/permissions`, `GET/POST /api/v1/memberships/:membershipId/roles`, `DELETE /api/v1/memberships/:membershipId/roles/:roleId`.

8.3 API keys

Purpose. Machine-to-machine credentials scoped to a tenant, used heavily by the compatibility surfaces.

A key has the wire format `prefix.secret`, where the prefix is `nxk_` plus eight alphanumeric characters and the secret is 32 random bytes encoded base64url. Only the **SHA-256 hash** of the secret is stored, alongside a globally unique prefix used for lookup. Verification hashes the presented secret and compares with a timing-safe comparison.

A key's scopes are intersected with the creating actor's own permissions at creation time, so a key can never be more privileged than the person who created it. At verification the principal's permissions are the stored scopes; live role membership is not re-resolved, which means revoking a user's role does not retroactively narrow a key they issued.

Rotation requires both the `api_keys.manage` permission **and** a valid MFA step-up session, and preserves the prefix while replacing the secret. Revocation requires the permission but not step-up.

Endpoints. `GET/POST /api/v1/api-keys`, `POST /api/v1/api-keys/:apiKeyId/rotate`, `POST /api/v1/api-keys/:apiKeyId/revoke`.

8.4 MFA

Purpose. Provide step-up assurance for sensitive operations.

TOTP secrets are generated with `otplib` and stored **encrypted with AES-256-GCM** using a key derived from `AUTH_MFA_ENCRYPTION_KEY`, which configuration validation requires to be exactly 32 bytes. Enrolment is a three-step flow — start, verify, activate — and activation issues recovery codes stored as SHA-256 hashes for one-time use.

A successful `POST /api/v1/mfa/verify` records a row in `step_up_sessions` valid for `AUTH_STEP_UP_TTL_SECONDS`. Operations such as API key rotation and webhook secret rotation check for that row.

Endpoints. `POST /api/v1/mfa/totp/start`, `/totp/verify`, `/totp/activate`, `POST /api/v1/mfa/verify`, `POST /api/v1/mfa/recovery-code/use`.

8.5 Tenants

Purpose. The tenancy root. Every tenant-scoped table keys off `tenants.id`.

Slugs are normalised to lowercase and must match `^[a-z0-9](?:[a-z0-9-]{0,61}[a-z0-9])?$`, enforced both in application code and by a database CHECK constraint. Status is `ACTIVE`, `SUSPENDED` or `CLOSED`; only `ACTIVE` tenants can authenticate.

Tenant isolation is enforced on two levels. The application calls `database.execute(..., { tenantId })`, which issues `set_config('app.tenant_id', <uuid>, true)` for the transaction. PostgreSQL row-level security policies then restrict rows to `tenant_id = app.current_tenant_id()`.

Endpoints. `POST /api/v1/tenants` and `GET /api/v1/tenants/:tenantId` are public by design (bootstrap and lookup); `suspend`, `reactivate` and `close` are protected.

8.6 Audit

Purpose. An append-only record of security-relevant and commerce-relevant actions.

Rows carry actor kind (`user`, `api-key`, `system`), actor id, event type, resource type and id, a JSONB metadata blob, IP address, request id and timestamp. Metadata passes through `redactAuditMetadata` before storage so tokens and secrets are not persisted. `AUDIT_EVENT_TYPES` in `domain/audit-event.js` enumerates 59 canonical event types.

Endpoint. `GET /api/v1/audit`, requiring `audit.read`, filterable by event type with a limit between 1 and 200.

8.7 Products

Purpose. The tenant catalogue.

A product carries `merchantSku` (unique per tenant, maximum 128 characters, trimmed but case-preserving), an optional `externalReference`, a product type and a status of `ACTIVE`, `INACTIVE` or `ARCHIVED`. Archiving is terminal: an archived product cannot be updated. Localised merchandising data lives separately in `product_content` keyed by locale, so adding a language does not touch the product row.

Endpoints. `GET/POST /api/v1/products`, `GET/PATCH /api/v1/products/:productId`, `POST .../deactivate`, `POST .../archive`, `GET .../content`, `PUT .../content/:locale`.

Events. `product.created`, `product.updated`, `product.status_changed`.

8.8 Offers

Purpose. The product-to-channel relationship. An offer is what makes a product sellable on a specific channel.

An offer is unique per tenant, product and channel. It holds the channel-side `externalReference`, a `priceReference`, an offer status of `DRAFT`, `ACTIVE`, `INACTIVE` or `SUSPENDED`, and a separate listing status. Activation can be configured to require that an effective price exists for the channel and currency, which prevents publishing an unpriced listing.

The offer's `externalReference` is the hinge for outbound synchronisation: catalog sync will skip an offer whose external reference is empty, because there is nothing on the marketplace side to update.

Endpoints. `GET/POST /api/v1/offers`, `GET/PATCH /api/v1/offers/:offerId`, `POST /api/v1/offers/:offerId/activate`.

8.9 Pricing

Purpose. Currency-correct prices with time validity.

A price row is `(product, optional channel, currency, amountMinor, validFrom, validTo, status)`. Amounts are integer minor units — there is no floating-point money anywhere in the domain. `validTo` must be strictly after `validFrom`, and inactive prices are immutable.

`getEffectivePrice` resolves the price applicable to a product on a channel at an instant. A channel-specific price takes precedence over a tenant-wide price, and among candidates the latest `validFrom` wins. Order creation and price synchronisation both call this single method, so an order and a marketplace listing cannot disagree about which price rule applied.

Endpoints. `GET/POST /api/v1/prices`, `GET/PATCH /api/v1/prices/:priceId`.

8.10 Inventory

Purpose. Authoritative stock, per product, per stock location.

The balance model has three fields with a hard invariant: `available = on_hand - reserved`, all non-negative, and `reserved ≤ on_hand`. Every operation both mutates the balance under a `SELECT ... FOR UPDATE` row lock and appends an immutable row to `inventory_movements`.

Operation	on_hand	reserved	available	Movement type
Reserve	—	increase	decrease	RESERVATION
Release	—	decrease	increase	—
Receive	increase	—	increase	RECEIPT
Adjust positive	increase	—	increase	ADJUSTMENT
Adjust negative	decrease	—	decrease	ADJUSTMENT
Fulfil reserved on ship	decrease	decrease	—	SALE
Record sale, no reservation	decrease	—	decrease	SALE
Record return	increase	—	increase	RETURN

The distinction between the last three rows matters operationally. A normal order reserves at creation and converts reservation to sale when the shipment ships, leaving `available` unchanged at that moment because it was already reduced at reservation. A **channel-fulfilled** order — where the marketplace shipped the goods itself — never reserved, so shipping records a direct sale and `available` drops then.

Endpoints. `GET /api/v1/inventory`, `GET /api/v1/inventory/:productId`, `POST /api/v1/inventory/adjustments`, `/receipts`, `/reservations`, `/releases`, plus `GET/POST /api/v1/stock-locations` and `GET /api/v1/stock-locations/:stockLocationId`.

Limitations. `TRANSFER_IN` and `TRANSFER_OUT` movement types exist in the enum but no endpoint or use case creates them. The release endpoint reuses the `inventory.reserve` permission rather than a distinct one.

8.11 Orders

Purpose. Order capture and fulfilment coordination — the busiest module in the system.

An order holds monetary totals in minor units, a tenant-unique `orderNumber` generated from `tenant_order_sequences`, an optional `externalOrderReference` for channel-originated orders, a customer snapshot, and lines. Each line tracks `cancelledQuantity`, `shippedQuantity` and `returnedQuantity`, from which the derived `cancellableQuantity`, `shippableQuantity` and `returnableQuantity` are computed. These derivations are why a partially shipped, partially cancelled order behaves correctly without extra bookkeeping.

There are four ways an order enters the system, and they differ in starting status:

Entry point	Starting status	Reserves stock
<code>POST /api/v1/orders</code> (native)	<code>CONFIRMED</code>	Yes
<code>POST /api/v2/orders</code> (channel ingest)	<code>NEW</code>	Yes
<code>POST /api/v2/orders/channel-fulfilled</code>	<code>NEW</code>	No — marketplace already shipped
Marketplace ingestion, polling or lifecycle	<code>NEW</code>	Yes

Endpoints. `GET/POST /api/v1/orders`, `GET /api/v1/orders/:orderId`, `POST /api/v1/orders/:orderId/confirm`. Channel creation and acknowledgement are exposed through `OrderCommandService` and the `/api/v2` routes rather than as native REST.

8.12 Shipments

Purpose. Represent a physical dispatch and trigger inventory consumption.

Statuses are `CREATED`, `READY_TO_SHIP`, `SHIPPED`, `IN_TRANSIT`, `DELIVERED`, `FAILED`, `CANCELLED`, with transitions asserted in `domain/shipment-status.js`. Creating a shipment validates each line against the order line's `shippableQuantity`. Shipping calls `fulfillInventoryWhenShipmentShipped`, which selects between `fulfillReservedForShipment` and `recordSale` based on the order's inventory consumption mode.

Endpoints. `POST /api/v1/orders/:orderId/shipments`, `GET /api/v1/shipments`, `GET /api/v1/shipments/:shipmentId`, `POST .../ship`, `/deliver`, `/cancel`.

8.13 Returns

Purpose. Handle goods coming back, with an approval gate before stock is restored.

The status flow is `REQUESTED` → `APPROVED` (or `REJECTED` / `CANCELLED`) → `RECEIVED` → `COMPLETED`. Stock is restored at **receive**, not at request, and receive is refused unless the return was approved. Receiving calls `inventoryService.recordReturn` and increments the order line's returned quantity.

Endpoints. `POST /api/v1/orders/:orderId/returns`, `GET /api/v1/returns`, `GET /api/v1/returns/:returnId`, `POST .../approve`, `/receive`, `/complete`, `/reject`, `/cancel`.

8.14 Cancellations

Purpose. Cancel order lines and release their reservations.

Unlike returns, cancellations complete synchronously: creating one releases inventory and marks the cancellation `COMPLETED` in a single transaction, emitting both `cancellation.created` and `cancellation.completed`. The `REJECTED` status exists on the entity but no HTTP path produces it.

Endpoints. `GET/POST /api/v1/cancellations`, `GET /api/v1/cancellations/:cancellationId`, `POST /api/v1/orders/:orderId/cancel`.

Note. `POST /api/v1/cancellations` requires an `Idempotency-Key`; `POST /api/v1/orders/:orderId/cancel` does not.

8.15 Channels and marketplaces

Channels are tenant-scoped sales destinations. A channel references a marketplace, optionally carries a `defaultStockLocationId` — the warehouse whose availability is published to that channel — and an `externalReference`. Marketplace credentials attach to the channel through a `marketplace_connections` row with encrypted credentials.

Marketplaces are the global registry plus all provider adapter code. Adapters are described in their own chapter.

Endpoints. `GET/POST /api/v1/channels`, `GET/PATCH /api/v1/channels/:channelId`, the four `.../marketplace-connection` verbs, `POST .../marketplace-connection/test`, and `GET/POST /api/v1/marketplaces` with `GET/PATCH /api/v1/marketplaces/:id`.

8.16 Webhooks

Purpose. Let a tenant subscribe to integration events and receive signed HTTP callbacks.

A subscription stores a destination URL, an encrypted secret and a list of event types validated against the catalogue's external-delivery allowlist. Delivery rows are unique per subscription and event, which is the dedupe mechanism. Delivery details are in the webhooks chapter.

Endpoints. `GET/POST /api/v1/webhooks`, `GET/PATCH/DELETE /api/v1/webhooks/:webhookId`, `GET .../deliveries`, `GET .../deliveries/:deliveryId`, `POST .../rotate-secret`.

8.17 External ID mapping

Purpose. ChannelEngine clients expect **integer** identifiers; Nexora uses UUIDs. This module bridges the two.

The `external_integer_id_mappings` table allocates a monotonically increasing integer per tenant, provider namespace and resource type. The namespace used by the compatibility layer is `compat_v2`, and the mapped resource types are `order`, `order_line`, `return`, `shipment` and `cancellation`. Mappings are assigned when the

entity is created and looked up when a compatibility response is built.

Products and channels are **not** in this table. Products are addressed by SKU, and channel integer identifiers are parsed from the channel's `externalReference` when that value happens to be numeric.

9 API architecture

9.1 Versioning and surfaces

Surface	Prefix	Contract owner	Notes
Native	/api/v1	Nexora	Full domain capability, UUID identifiers, { success, data } envelope
Merchant compatibility	/api/v2	ChannelEngine-shaped	PascalCase fields, integer ids, Idempotency-Key required on mutations
StockConnect CE	/api/v2/ce	ChannelEngine-shaped	Additive to /api/v2, extra credential channels, optional idempotency

ADR-010 records this strategy: /api/v1 evolves with Nexora's own needs while /api/v2 is a facade constrained by an external contract. The two never share request or response schemas; the compatibility layer maps between them explicitly.

9.2 Validation and serialisation

Every route declares a Zod schema. fastify-type-provider-zod installs Zod as both the validator compiler and the serializer compiler, so a request that does not match is rejected before the handler runs, and a response that does not match the declared shape fails serialisation. The same schemas feed @fastify/swagger, which is why the OpenAPI document is always consistent with enforced behaviour.

9.3 Success and error envelopes

Native routes return { success: true, data: ... }. Errors anywhere on the native surface return a single normalised envelope produced by src/app/errors/error-envelope.js:

```
{
  "success": false,
  "error": {
    "code": "VALIDATION_FAILED",
    "message": "Request validation failed",
    "details": { }
  },
  "requestId": "6f0f1e3a-..."
}
```

details is omitted when there is nothing safe to expose. When an error is marked non-exposable, message is replaced by the generic string An unexpected error occurred so internal details never leak. requestId is always present and is the value returned in the x-request-id response header, which makes support triage straightforward.

The compatibility surfaces use a different envelope — `{ Success, StatusCode, Message, ValidationErrors }` — because that is what ChannelEngine clients parse. This is produced by `mapCoreErrorToExternalApiResponse` and is described in the compatibility chapter.

9.4 Error taxonomy

Error class	HTTP	Code
<code>ValidationError</code>	400	<code>VALIDATION_FAILED</code>
<code>MalformedRequestError</code>	400	<code>MALFORMED_REQUEST</code>
<code>AuthenticationError</code>	401	<code>AUTHENTICATION_REQUIRED</code>
<code>AuthorizationError</code>	403	<code>FORBIDDEN</code>
<code>NotFoundError</code>	404	<code>NOT_FOUND</code>
<code>ConflictError</code>	409	<code>CONFLICT</code>
<code>IdempotencyConflictError</code>	409	<code>IDEMPOTENCY_CONFLICT</code>
<code>IdempotentRequestInProgressError</code>	409	<code>IDEMPOTENT_REQUEST_IN_PROGRESS</code>
<code>BusinessRuleError</code>	422	<code>BUSINESS_RULE_VIOLATION</code>
<code>RateLimitError</code>	429	<code>RATE_LIMIT_EXCEEDED</code>
<code>ExternalServiceError</code>	502	<code>EXTERNAL_SERVICE_ERROR</code>
<code>DatabaseError</code>	503	<code>DATABASE_UNAVAILABLE</code>
<code>CircuitOpenError</code>	503	<code>CIRCUIT_OPEN</code>
<code>ServiceUnavailableError</code>	503	<code>SERVICE_UNAVAILABLE</code>
<code>ExternalServiceTimeoutError</code>	504	<code>EXTERNAL_SERVICE_TIMEOUT</code>
<code>ConfigurationError</code>	500	<code>CONFIGURATION_INVALID</code>
<code>InternalError</code>	500	<code>INTERNAL_ERROR</code>

`RateLimitError` additionally sets a `Retry-After` header. The registry also defines `PAYLOAD_TOO_LARGE` and `UNSUPPORTED_MEDIA_TYPE`, but no mapping currently produces them, so an oversized body surfaces as a generic internal error.

9.5 Idempotency

Mutating endpoints on `/api/v2` require an `Idempotency-Key` header. The record is stored in `idempotency_records`, keyed by tenant, a principal fingerprint, a route identifier and the key. A stable SHA-256 fingerprint of the request body (`src/shared/idempotency/fingerprint.js`) detects the case where the same key

is reused with a different payload, which returns 409.

Separately, the commerce modules implement **request equivalence**: when a create request carries a natural business key that already exists — a channel order's external reference, a shipment's `externalReference` — the handler compares the new payload against the stored aggregate. If they are equivalent the existing record is returned; if they differ, the request is rejected as a conflict. This protects against duplicate ingestion even when the caller did not send an idempotency key.

9.6 Pagination, filtering and rate limiting

Native list endpoints use cursor pagination with limits clamped between 1 and 100

(`src/shared/pagination/cursor.js`). The compatibility surfaces use page-and-size pagination because that is what the external contract specifies — PascalCase `Page / PageSize` on `/api/v2` and lowercase `page / pageSize` on `/api/v2/ce` .

Rate limiting is implemented as a Redis GCRA algorithm in a Lua script (`src/infrastructure/redis/redis-rate-limiter.js`). There is no in-memory fallback: if Redis is unavailable, rate-limited operations fail closed. Named policies live in `src/shared/auth/rate-limit-policies.js` .

Policy	Limit	Window
<code>auth.login</code>	10	900 s
<code>auth.refresh</code>	30	900 s
<code>catalog-sync</code>	120	60 s
Compatibility read	120	60 s
Compatibility mutation	60	60 s

Compatibility rate limits are keyed per tenant **and** per credential, using a subject of the form `<tenantId>:api-key:<id>:read` , so one noisy API key cannot exhaust another's budget.

9.7 API documentation endpoints

When `DOCS_ENABLED` is true:

URL	Content
<code>/docs</code>	Scalar interactive API reference
<code>/api-docs</code>	Redirect to <code>/docs</code>
<code>/openapi.json</code>	OpenAPI 3.1 document
<code>/api-docs.json</code>	Same document under an alias path

10 API reference

The complete machine-readable reference is the generated OpenAPI document. This chapter gives the registered surface grouped by module, with the route identifiers used by the repository's route inventory so that any endpoint can be cross-referenced with `docs/API_ROUTE_INVENTORY.md`, `docs/SWAGGER_COVERAGE_REPORT.md` and the Postman collection.

10.1 Surface totals

Metric	Count
Unique registered endpoints on the API process	147
GET	58
POST	68
PUT	9
PATCH	9
DELETE	3
Worker observability endpoints (separate port)	3

10.2 Operational and documentation endpoints

Route ID	Method	Path	Auth	Purpose
R-0144	GET	/health/live	Public	Liveness; always { status: 'ok' } on the API process
R-0145	GET	/health/ready	Public	Runs all readiness probes; 503 when any fails
R-0146	GET	/internal/metrics	Public	Prometheus exposition
R-0147	GET	/openapi.json	Public when docs enabled	OpenAPI 3.1 document
R-0001	GET	/api-docs	Public when docs enabled	Redirect to Scalar UI
R-0002	GET	/api-docs.json	Public when docs enabled	OpenAPI alias
R-0024	POST	/api/v1/foundation/echo	Public	Validated echo for connectivity testing
R-0025	GET	/api/v1/foundation/ping	Public	Returns pong

`/internal/metrics` and `/health/*` are unauthenticated by design and must not be exposed publicly at the ingress.

10.3 Identity, tenancy and access control

Route ID	Method	Path	Auth	Purpose
R-0008	POST	<code>/api/v1/auth/login</code>	Public	Authenticate with tenant slug, email, password
R-0011	POST	<code>/api/v1/auth/refresh</code>	Public	Rotate refresh token, issue new access token
R-0009	POST	<code>/api/v1/auth/logout</code>	Public	Revoke the refresh session
R-0010	GET	<code>/api/v1/auth/me</code>	Bearer	Current principal and membership status
R-0087	POST	<code>/api/v1/tenants</code>	Public	Create a tenant
R-0088	GET	<code>/api/v1/tenants/:tenantId</code>	Public	Read a tenant
R-0091	POST	<code>/api/v1/tenants/:tenantId/suspend</code>	Protected	Suspend a tenant
R-0090	POST	<code>/api/v1/tenants/:tenantId/reactivate</code>	Protected	Reactivate a tenant
R-0089	POST	<code>/api/v1/tenants/:tenantId/close</code>	Protected	Close a tenant
R-0077	GET	<code>/api/v1/roles</code>	Protected	List tenant roles
R-0078	POST	<code>/api/v1/roles</code>	Protected	Create a role
R-0057	GET	<code>/api/v1/permissions</code>	Protected	List the permission catalogue
R-0037	GET	<code>/api/v1/memberships/:membershipId/roles</code>	Protected	Roles on a membership
R-0038	POST	<code>/api/v1/memberships/:membershipId/roles</code>	Protected	Assign a role
R-0039	DELETE	<code>/api/v1/memberships/:membershipId/roles/:roleId</code>	Protected	Remove a role, subject to the last-admin guard
R-0003	GET	<code>/api/v1/api-keys</code>	Protected	List keys; requires <code>api_keys.read</code>
R-0004	POST	<code>/api/v1/api-keys</code>	Protected	Create a key; secret returned once
R-0006	POST	<code>/api/v1/api-keys/:apiKeyId/rotate</code>	Protected + step-up	Rotate the secret
R-0005	POST	<code>/api/v1/api-keys/:apiKeyId/revoke</code>	Protected	Revoke a key
R-0042	POST	<code>/api/v1/mfa/totp/start</code>	Protected	Begin TOTP enrolment

Route ID	Method	Path	Auth	Purpose
R-0043	POST	/api/v1/mfa/totp/verify	Protected	Verify the enrolment code
R-0041	POST	/api/v1/mfa/totp/activate	Protected	Activate the factor, issue recovery codes
R-0044	POST	/api/v1/mfa/verify	Protected	Create a step-up session
R-0040	POST	/api/v1/mfa/recovery-code/use	Protected	Step up with a one-time recovery code
R-0007	GET	/api/v1/audit	Protected	Query the audit log; requires <code>audit.read</code>

10.4 Catalogue, pricing and inventory

Route ID	Method	Path	Purpose
R-0062 / R-0063	GET / POST	/api/v1/products	List and create products
R-0064 / R-0065	GET / PATCH	/api/v1/products/:productId	Read and update
R-0069	POST	/api/v1/products/:productId/deactivate	Set status <code>INACTIVE</code>
R-0066	POST	/api/v1/products/:productId/archive	Terminal archive
R-0067	GET	/api/v1/products/:productId/content	Localised content
R-0068	PUT	/api/v1/products/:productId/content/:locale	Upsert one locale
R-0045 / R-0046	GET / POST	/api/v1/offers	List and create offers
R-0047 / R-0048	GET / PATCH	/api/v1/offers/:offerId	Read and update
R-0049	POST	/api/v1/offers/:offerId/activate	Activate, optionally requiring an effective price
R-0058 / R-0059	GET / POST	/api/v1/prices	List and create prices
R-0060 / R-0061	GET / PATCH	/api/v1/prices/:priceId	Read and update
R-0084 / R-0085	GET / POST	/api/v1/stock-locations	List and create locations
R-0086	GET	/api/v1/stock-locations/:stockLocationId	Read a location
R-0027	GET	/api/v1/inventory	Balances across locations

Route ID	Method	Path	Purpose
R-0028	GET	/api/v1/inventory/:productId	Balances for one product
R-0029	POST	/api/v1/inventory/adjustments	Signed adjustment
R-0030	POST	/api/v1/inventory/receipts	Receive stock
R-0032	POST	/api/v1/inventory/reservations	Reserve stock
R-0031	POST	/api/v1/inventory/releases	Release a reservation

10.5 Orders, shipments, returns and cancellations

Route ID	Method	Path	Purpose
R-0050	GET	/api/v1/orders	List orders
R-0051	POST	/api/v1/orders	Create a native order; Idempotency-Key required
R-0052	GET	/api/v1/orders/:orderId	Read an order
R-0054	POST	/api/v1/orders/:orderId/confirm	NEW → CONFIRMED
R-0056	POST	/api/v1/orders/:orderId/shipments	Create a shipment; Idempotency-Key required
R-0079 / R-0080	GET	/api/v1/shipments, /:shipmentId	List and read
R-0083	POST	/api/v1/shipments/:shipmentId/ship	Ship and consume inventory
R-0082	POST	/api/v1/shipments/:shipmentId/deliver	Mark delivered
R-0081	POST	/api/v1/shipments/:shipmentId/cancel	Cancel and reverse quantities
R-0055	POST	/api/v1/orders/:orderId/returns	Create a return; Idempotency-Key required
R-0070 / R-0071	GET	/api/v1/returns, /:returnId	List and read
R-0072	POST	/api/v1/returns/:returnId/approve	Approve
R-0075	POST	/api/v1/returns/:returnId/receive	Receive goods and restock
R-0074	POST	/api/v1/returns/:returnId/complete	Complete
R-0076 / R-0073	POST	/api/v1/returns/:returnId/reject, /cancel	Terminal outcomes
R-0012 / R-0013	GET / POST	/api/v1/cancellations	List and create; create requires Idempotency-Key
R-0014	GET	/api/v1/cancellations/:cancellationId	Read

Route ID	Method	Path	Purpose
R-0053	POST	<code>/api/v1/orders/:orderId/cancel</code>	Cancel by order; no idempotency wrapper

10.6 Channels, marketplaces and webhooks

Route ID	Method	Path	Purpose
R-0015 / R-0016	GET / POST	<code>/api/v1/channels</code>	List and create channels
R-0017 / R-0018	GET / PATCH	<code>/api/v1/channels/:channelId</code>	Read and update
R-0022	POST	<code>/api/v1/channels/:channelId/marketplace-connection</code>	Create the connection with encrypted credentials
R-0020	GET	<code>.../marketplace-connection</code>	Read the redacted connection
R-0021	PATCH	<code>.../marketplace-connection</code>	Update
R-0019	DELETE	<code>.../marketplace-connection</code>	Remove
R-0023	POST	<code>.../marketplace-connection/test</code>	Live credential test with sanitised errors
R-0033 / R-0034	GET / POST	<code>/api/v1/marketplaces</code>	Registry list and create
R-0035 / R-0036	GET / PATCH	<code>/api/v1/marketplaces/:id</code>	Read and update
R-0092 / R-0093	GET / POST	<code>/api/v1/webhooks</code>	List and create subscriptions
R-0095 / R-0096 / R-0094	GET / PATCH / DELETE	<code>/api/v1/webhooks/:webhookId</code>	Manage a subscription
R-0097	GET	<code>/api/v1/webhooks/:webhookId/deliveries</code>	Delivery history
R-0098	GET	<code>.../deliveries/:deliveryId</code>	One delivery attempt record
R-0099	POST	<code>/api/v1/webhooks/:webhookId/rotate-secret</code>	Rotate the signing secret; requires step-up
R-0026	POST	<code>/api/v1/inbound/marketplace-webhooks/:ingressToken</code>	Inbound marketplace webhook; authenticated by the ingress token

10.7 Compatibility surfaces

The `/api/v2` and `/api/v2/ce` endpoints (route identifiers R-0100 through R-0143) are documented endpoint by endpoint in the ChannelEngine compatibility chapter, because their semantics are only meaningful alongside the mapping and idempotency rules described there.

10.8 Worker observability endpoints

These are served by the worker process on `WORKER_OBSERVABILITY_PORT`, default 3001, bound to `127.0.0.1` by default. They are **not** reachable through the API base URL.

Method	Path	Purpose
GET	<code>/health/live</code>	Returns 503 with <code>{ status: 'shutting_down' }</code> during drain
GET	<code>/health/ready</code>	Worker probe set: postgres, redis, queue, workers registered
GET	<code>/internal/metrics</code>	Prometheus exposition for the worker process

11 Authentication and authorization

11.1 Credential types

Credential	Transport	Resolved by	Principal kind
JWT access token	Authorization: Bearer <jwt>	authenticateAccessToken	user
API key	x-api-key: <key> or Authorization: ApiKey <key>	verifyApiKey	api-key
CE API key	Query apiKey or apikey, or header X-CE-KEY	verifyApiKey, only on /api/v2/ce/*	api-key
Webhook ingress token	Path segment on the inbound webhook route	Hash lookup against the connection	none — resolved to a connection

Supplying both a bearer token and an API key on the same request is an error: the plugin raises `AuthenticationError('Provide either Bearer token or API key, not both')`.

11.2 The public route allowlist


Authentication is skipped only for paths matching this list, taken verbatim from

`src/app/http/plugins/authentication.plugin.js`:

```
const PUBLIC_ROUTE_PATTERNS = [  
  /^\/health\/$/,  
  /^\/internal\/metrics$/,  
  /^\/docs(?:\/|$)/$,  
  /^\/openapi\.json$/,  
  /^\/api-docs\.json$/,  
  /^\/api-docs$/,  
  /^\/api\/v1\/foundation\/$/,  
  /^\/api\/v1\/auth\/login$/,  
  /^\/api\/v1\/auth\/refresh$/,  
  /^\/api\/v1\/auth\/logout$/,  
  /^\/api\/v1\/tenants$/,  
  /^\/api\/v1\/inbound\/marketplace-webhooks\/.+$/,  
];
```

Two additional cases are handled in code rather than by regex: `POST /api/v1/tenants` and any `GET /api/v1/tenants/...`.

11.3 Authorization flow



There are two independent enforcement layers. Permission checks decide *what action* is allowed; row-level security decides *which rows* are visible. A bug in one does not automatically defeat the other.

11.4 Secret storage

Secret	Algorithm
User password	Argon2id, library default parameters
Refresh token	SHA-256 hex of the raw token
Password reset token	SHA-256 hex
MFA recovery code	SHA-256 hex
API key secret	SHA-256 hex, timing-safe comparison
MFA TOTP secret	AES-256-GCM at rest
Webhook secret, marketplace credentials	Encrypted ciphertext columns
JWT	HS256 with a shared secret, or RS256 with a key pair

11.5 Permission catalogue

Permission keys follow `^[a-z][a-z0-9_]*$`. Those present in migrations and code include:

tenant.admin, orders.read, orders.create, orders.update, orders.cancel, orders.ingest, orders.ingest_channel_fulfilled, products.read, products.create, products.update, inventory.read, inventory.update, inventory.adjust, inventory.reserve, shipments.read, shipments.create, shipments.update, returns.read, returns.create, returns.update, cancellations.read, cancellations.create, channels.read, channels.create, channels.update, users.read, users.manage, roles.read, roles.manage, audit.read, audit.export, api_keys.read, api_keys.manage, mfa.manage, auth.step_up, marketplaces.read, marketplaces.manage, pricing.read, pricing.create, pricing.update, offers.read, offers.create, offers.update, webhooks.read, webhooks.manage.

ADR-012 defines seven seeded system roles per tenant, expanded from these keys through the pattern mechanism.

11.6 Verified security limitations

These are stated neutrally as facts about the current code, not as vulnerabilities with assessed severity.

1. **Unscoped transactions on RLS-protected tables.** Refresh, logout and API-key verification query `refresh_sessions` and `api_keys` through `database.execute()` **without** a `tenantId`, because the tenant is not yet known. With `app.current_tenant_id()` unset, an RLS policy would deny the rows. These paths therefore depend on the connecting role bypassing RLS, which is inconsistent with the `nexora_app` restricted-role model introduced by migration `0030_rls_hardening.sql`.
2. **Hardcoded role password in a migration.** `0030_rls_hardening.sql` creates the `nexora_app` role with a literal password, documented as a Docker development convenience.
3. **Default MFA encryption key.** `AUTH_MFA_ENCRYPTION_KEY` has a default value in the configuration schema; deployments that do not override it encrypt TOTP secrets with a known key.
4. **No MFA at login.** MFA gates step-up operations only.
5. **Exact-match authorization only at runtime.** Wildcard permission patterns apply at role seeding, not at request time.
6. **API key permissions are frozen at issuance.** Scopes are stored on the key; later role changes do not narrow them.
7. **Content Security Policy disabled.** Helmet is registered with `contentSecurityPolicy: false`.
8. **Broad unauthenticated surface by design.** Tenant creation and lookup, health, metrics, documentation and marketplace webhook ingress are all public.
9. **SECURITY DEFINER webhook lookup.** Migration `0047` adds `app.lookup_marketplace_connection_for_webhook`, which intentionally bypasses the tenant GUC so an inbound webhook can be routed before the tenant is known.
10. **Global tables without RLS.** `users`, `permissions` and the password tables are global; access control for them is purely application-level.

12 Database architecture

12.1 Technology and access pattern

PostgreSQL is accessed through `pg` with hand-written SQL in repository classes. There is no ORM, a deliberate choice recorded in ADR-011. Pool behaviour, statement timeouts and query timeouts are all configurable, and every connection sets `application_name` to `nexora-backend` so sessions are identifiable in `pg_stat_activity`.

Transactions are opened by the application layer, optionally with an isolation level and read-only flag, and optionally with a tenant identifier that sets the transaction-local GUC used by row-level security.


12.2 Migrations

Migrations are forward-only SQL files applied by a custom runner (`src/infrastructure/postgres/migrator.js`). Each file's checksum is recorded, so editing an applied migration is detected. The runner takes PostgreSQL advisory lock `8147302915` to make concurrent startups safe, and records applied versions in `schema_migrations`.


Forty-seven migrations exist, in order:

Range	Theme
0001 – 0004	App schema, tenant context function, outbox, inbox, idempotency
0005 – 0011	Tenants, users, memberships, roles, permissions, authentication, API keys, MFA, audit
0012 – 0017	Marketplaces, channels, products, inventory, pricing, offers
0018 – 0022	Commerce permissions, composite foreign keys, permission backfills
0023 – 0026	Orders, shipments, cancellations, returns
0027 – 0030	Phase 4 permissions and composite keys, RLS hardening
0031 – 0033	External reference columns on shipments, cancellations, returns
0034 – 0037	Webhook subscriptions, deliveries, permissions
0038 – 0041	Order ingestion permissions and backfills
0042 – 0044	External integer id mappings, channel default stock location, delivery retention
0045 – 0047	Marketplace connections, entity mapping, webhook ingress

12.3 Core entity relationships



12.4 Identity and access tables



12.5 Table inventory

Table	Purpose	Tenant-scoped	RLS
<code>schema_migrations</code>	Applied migration ledger	No	No
<code>tenants</code>	Tenant registry with unique slug	—	No
<code>users</code>	Global user identity, unique normalised email	No	No
<code>password_credentials</code>	Argon2id password hashes	No	No
<code>password_reset_tokens</code>	Hashed reset tokens	No	No
<code>tenant_memberships</code>	User-to-tenant link with status	Yes	Yes
<code>permissions</code>	Global permission catalogue	No	No
<code>roles</code>	Tenant roles, system and custom	Yes	Yes

Table	Purpose	Tenant-scoped	RLS
role_permissions	Role to permission grants	Yes	Yes
membership_roles	Membership to role assignment	Yes	Yes
refresh_sessions	Rotating refresh tokens with family tracking	Yes	Yes
api_keys	Hashed API key secrets and scopes	Yes	Yes
mfa_factors	AES-256-GCM encrypted TOTP secrets	Yes	Yes
recovery_codes	Hashed one-time recovery codes	Yes	Yes
step_up_sessions	Step-up assurance records	Yes	Yes
audit_log	Append-only audit trail	Nullable	Yes
outbox_events	Transactional outbox	Optional	No
inbox_messages	Consumer deduplication	No	No
idempotency_records	HTTP idempotency ledger	Optional	No
marketplaces	Global marketplace definitions	No	No
channels	Tenant sales channels	Yes	Yes
marketplace_connections	Encrypted marketplace credentials and webhook ingress token hash	Yes	Yes
marketplace_entity_mappings	Nexora entity to external entity identifiers	Yes	Yes
products	Tenant products, unique merchant SKU	Yes	Yes
product_content	Localised product content	Yes	Yes
stock_locations	Warehouses	Yes	Yes
inventory_balances	On hand, reserved, available	Yes	Yes
inventory_movements	Append-only stock ledger	Yes	Yes
inventory_reservations	Active and released reservations	Yes	Yes
prices	Minor-unit prices with validity windows	Yes	Yes
offers	Product-channel listings	Yes	Yes
tenant_order_sequences	Per-tenant order number counter	Yes	No
orders	Order header with totals and status	Yes	Yes
order_lines	Lines with cancelled, shipped, returned counters	Yes	Yes

Table	Purpose	Tenant-scoped	RLS
<code>order_customer_snapshots</code>	Customer PII captured at order time	Yes	Yes
<code>shipments / shipment_lines</code>	Shipment documents and allocations	Yes	Yes
<code>returns / return_lines</code>	Return requests and lines	Yes	Yes
<code>cancellations / cancellation_lines</code>	Cancellations and lines	Yes	Yes
<code>webhook_subscriptions</code>	Destination, encrypted secret, event types	Yes	Yes
<code>webhook_deliveries</code>	Delivery attempts, unique per subscription and event	Yes	Yes
<code>external_integer_id_sequences</code>	Compatibility integer counters	Yes	No
<code>external_integer_id_mappings</code>	UUID to integer mapping for <code>compat_v2</code>	Yes	Yes

12.6 Tenant isolation mechanics

Three mechanisms work together:

1. **Transaction-local GUC.** `set_config('app.tenant_id', <uuid>, true)` scopes the setting to the transaction, so a pooled connection cannot leak a tenant context to the next checkout.
2. **Row-level security policies.** Tenant-scoped tables restrict rows to `tenant_id = app.current_tenant_id()`. The audit log additionally permits `tenant_id IS NULL` for system events.
3. **Composite foreign keys.** Migrations `0019`, `0020`, `0028` and `0046` add `(tenant_id, id)` composite keys so that a child row cannot reference a parent belonging to a different tenant, even if application code were to try.

Migration `0030_rls_hardening.sql` adds `FORCE ROW LEVEL SECURITY` and the dedicated `nexora_app` role, with the expectation that migrations run under a privileged role via `DATABASE_MIGRATION_URL` while runtime traffic uses the restricted role.

12.7 Data retention

A scheduled cleanup running on the worker deletes aged rows in bounded batches until fewer than a full batch remain.

Data	Environment variable	Default
Published outbox events	<code>OUTBOX_RETENTION_DAYS</code>	30
Processed inbox messages	<code>INBOX_RETENTION_DAYS</code>	30
Expired idempotency records	<code>IDEMPOTENCY_RETENTION_DAYS</code>	7
Terminal webhook deliveries	<code>WEBHOOK_DELIVERY_RETENTION_DAYS</code>	30

Data	Environment variable	Default
Batch size	RETENTION_CLEANUP_BATCH_SIZE	100
Interval	RETENTION_CLEANUP_INTERVAL_MS	3600000

The job holds a distributed lock named `retention-cleanup`, so running several worker replicas does not cause concurrent deletion. Per-resource failures are recorded individually and exported through `nexora_retention_cleanup_*` metrics.

13 ChannelEngine integration for StockConnect

This chapter is the core of the document. It describes what the ChannelEngine relationship actually is in this codebase, how the compatibility surface is built, how each data domain is mapped, and where the implementation is complete, partial or absent.

13.1 What the integration actually is

StockConnect is an existing external merchant system that was written against **ChannelEngine's** API. ChannelEngine is a third-party channel-management SaaS. The business goal recorded in the planning material is to **replace ChannelEngine with Nexora** in StockConnect's architecture without rewriting StockConnect.

The technique chosen is a **compatibility facade**. Nexora implements HTTP endpoints whose paths, field names, pagination style, identifier types and error envelopes imitate ChannelEngine closely enough that StockConnect can change its base URL and credentials and keep working.


This produces a relationship that is the opposite of what the phrase "ChannelEngine integration" normally implies:

Common assumption	Actual implementation
Nexora calls ChannelEngine's API	Nexora exposes a ChannelEngine-shaped API and StockConnect calls it
There is a ChannelEngine API client	There is none; no outbound HTTP to ChannelEngine exists
ChannelEngine credentials are configured	No <code>CHANNELENGINE_*</code> environment variables exist
ChannelEngine reaches the marketplaces	Nexora's own marketplace adapters do that, independently

The evidence is direct. Searching `src/` for "ChannelEngine" returns three files, all within the compatibility module, and every occurrence is a comment or an OpenAPI description.

`src/modules/compatibility/infrastructure/index.js` is an empty boundary stub — the module has no infrastructure because it makes no outbound calls. Its `domain/index.js` is likewise an empty placeholder, because the compatibility layer holds no business rules of its own; it translates and delegates.

13.2 System context




The dotted line is the point of the diagram: ChannelEngine is present in the picture only as the system being replaced.

13.3 Responsibility split

Responsibility	Owner
Calling the compatibility API, polling for orders	StockConnect
Presenting a ChannelEngine-shaped contract	Nexora compatibility module
Authenticating the caller and enforcing tenancy	Nexora authentication plugin and RLS
Validating and mapping payloads	Nexora compatibility mappers
Executing business logic	Nexora core commerce modules
Allocating stable integer identifiers	Nexora external-id-mapping module
Reaching Amazon, Noon, Namshi, Shopify	Nexora marketplace adapters
Anything ChannelEngine used to do beyond these	Out of scope; not reimplemented

13.4 Compatibility module architecture



Both route plugins install their **own** Fastify error handler that converts errors to the CE envelope, so a compatibility client never sees Nexora's native error shape.

13.5 Authentication for compatibility clients

Core authentication accepts a bearer JWT or an API key via `x-api-key` or `Authorization: ApiKey <key>`. The `/api/v2/ce/*` prefix accepts two **additional** credential channels, because ChannelEngine clients conventionally pass the key in the query string:

```

if (apiKey === null && isMerchantCompatQueryAuthPath(path)) {
  apiKey = readMerchantCompatQueryApiKey(request.query)
  ?? readMerchantCompatCeKeyHeader(request.headers);
}

```

Channel	Form	Accepted on
Bearer JWT	Authorization: Bearer <jwt>	All protected routes
API key header	x-api-key: <key>	All protected routes
API key scheme	Authorization: ApiKey <key>	All protected routes
Query parameter	?apiKey=<key> or ?apikey=<key>	/api/v2/ce/* only
CE header	X-CE-KEY: <key>	/api/v2/ce/* only

The prefix restriction is enforced by `isMerchantCompatQueryAuthPath` in `src/shared/auth/merchant-compat-route-prefix.js`. An integration test confirms the boundary: a query `apiKey` on `GET /api/v2/orders` returns 401, while the same credential on a `/api/v2/ce` route succeeds. All five channels resolve to the same `verifyApiKey` verification path and therefore the same tenant scoping and RLS behaviour.

Operational note. Query-string credentials appear in access logs and proxy logs at any intermediary. The logging redaction configuration covers `x-api-key` and `x-ce-key` in Nexora's own logs, but an upstream load balancer will still record the full URL. This is an accepted consequence of the ChannelEngine contract, not an oversight.

13.6 Endpoint inventory

Merchant surface, `/api/v2/*`

Route ID	Method	Path	Purpose	Idempotency	Delegates to
R-0122	GET	/api/v2/foundation/ping	Connectivity probe	—	Inline
R-0125	GET	/api/v2/orders	List orders with filters	—	<code>OrderCompatibilityQuery.listOrders</code>
R-0129	GET	/api/v2/orders/new	Orders in <code>NEW</code> status	—	<code>listNewOrders</code>
R-0126	POST	/api/v2/orders	Channel order ingest, reserves stock	Required	<code>createChannelOrder</code>
R-0128	POST	/api/v2/orders/channel-fulfilled	Ingest an order the channel already shipped	Required	<code>createChannelFulfilledOrder</code>
R-0127	POST	/api/v2/orders/acknowledge	Acknowledge import	Required	<code>acknowledgeOrder</code>

Route ID	Method	Path	Purpose	Idempotency	Delegates to
R-0143	GET	/api/v2/shipments/merchant	List shipments	—	listMerchantShipments
R-0141	POST	/api/v2/shipments	Create a shipment	Required	createShipment
R-0142	PUT	/api/v2/shipments/:merchantShipmentNo	Update tracking, transition to shipped	Required	updateShipmentTracking
R-0101	GET	/api/v2/cancellations/merchant	List cancellations	—	listMerchantCancellations
R-0100	POST	/api/v2/cancellations	Create a cancellation	Required	createCancellation
R-0140	GET	/api/v2/returns/merchant/new	Unhandled returns	—	listNewMerchantReturns
R-0138	GET	/api/v2/returns/merchant/:merchantOrderNo	Returns for one order	—	listReturnsByMerchantOrderNo
R-0137	GET	/api/v2/returns/merchant	List returns	—	listMerchantReturns
R-0135	POST	/api/v2/returns	Create a return	Required	createReturn
R-0136	PUT	/api/v2/returns	Receive, accept or reject	Required	receiveReturn
R-0139	POST	/api/v2/returns/merchant/acknowledge	Acknowledge a return	Required	acknowledgeReturn
R-0130	GET	/api/v2/products	Products by merchant product numbers	—	listProductsByMerchantProductNos
R-0131	POST	/api/v2/products	Batch product upsert	Required	upsertProducts
R-0134	POST	/api/v2/products/freeze	Suspend offers or deactivate	Required	freezeProducts
R-0132	POST	/api/v2/products/bulkdelete	Soft deactivate	Required	bulkDeleteProducts
R-0133	PATCH	/api/v2/products/extra-data/bulk	Merge CE ExtraData	Required	patchExtraDataBulk
R-0123	PUT	/api/v2/offer	Update offer price	Required	updateOfferPrice
R-0124	PUT	/api/v2/offer/stock	Set absolute stock	Required	updateOfferStock

StockConnect CE surface, /api/v2/ce/*


Route ID	Method	Path	Purpose	Idempotency	Delegates to
R-0107	GET	/api/v2/ce/orders	Order poll, page and pageSize	—	listOrdersForStockConnectPoll

Route ID	Method	Path	Purpose	Idempotency	Delegates to
R-0108	GET	/api/v2/ce/orders/:merchantOrderNo/invoice	Invoice PDF	—	StockConnectCeOrderInvoiceQuery
R-0109	POST	/api/v2/ce/orders/acknowledge	Acknowledge	Optional, derived	acknowledgeOrder
R-0102	POST	/api/v2/ce/cancellations	Create a cancellation	Optional, derived	createCancellation
R-0119	POST	/api/v2/ce/shipments	Create a shipment	Optional, derived	createShipment
R-0121	GET	/api/v2/ce/shipments/merchant	List shipments	—	listMerchantShipments
R-0120	PUT	/api/v2/ce/shipments/:merchantShipmentNo/delivery-state	Report delivery state	Computed, not persisted	StockConnectCeShipmentDeliveryCommand
R-0115	GET	/api/v2/ce/returns	Return poll	—	listMerchantReturns
R-0117	POST	/api/v2/ce/returns/merchant	Create a return	Optional, derived	createReturn
R-0118	POST	/api/v2/ce/returns/merchant/acknowledge	Acknowledge a return	Optional, derived	acknowledgeReturn
R-0116	PUT	/api/v2/ce/returns	Receive a return	Optional, derived	receiveReturn
R-0110	GET	/api/v2/ce/products	Products by SKU list	—	StockConnectCeProductsQuery
R-0103	GET	/api/v2/ce/channels	Channel registry	—	StockConnectCeChannelCompatibilityQuery
R-0104	GET	/api/v2/ce/channels/:channelId/products	Channel listing status	—	StockConnectCeChannelProductsQuery
R-0111	POST	/api/v2/ce/products	Product push	Computed, not persisted	StockConnectCeCatalogCommand.pushProducts
R-0106	PUT	/api/v2/ce/offer/stock	Set stock	Computed, not persisted	updateOfferStock
R-0105	PUT	/api/v2/ce/offer	Set price	Computed, not persisted	updateOfferPrice
R-0114	POST	/api/v2/ce/products/freeze	Freeze or unfreeze	Computed, not persisted	freezeProducts
R-0112	POST	/api/v2/ce/products/bulkdelete	Bulk delete or archive	Computed, not persisted	bulkDeleteProducts
R-0113	PATCH	/api/v2/ce/products/extra-data/bulk	Extra data merge	Computed, not persisted	patchExtraData

Routes present on `/api/v2` but deliberately **absent** from `/api/v2/ce`: `foundation/ping`, `orders/new`, filtered order listing, channel order ingest by POST, shipment tracking update by PUT, cancellation listing, and the per-order and new-return sub-routes.

13.7 Identifier mapping

ChannelEngine clients expect integer identifiers. Nexora uses UUIDs internally. The bridge is the `external-id-mapping` module.



Integers are allocated monotonically per tenant, provider and resource type at the moment the entity is created — `assignOrderCompatibilityExternalIds` and its siblings run inside the creation transaction. Mapped resource types are `order`, `order_line`, `return`, `shipment` and `cancellation`.

Inbound resolution is handled by `compatibility-external-id-resolution.js`, which is strict about consistency. `resolveOrderForCompatibility` resolves `MerchantOrderNo` to an order number, and if the caller **also** supplies an `OrderId`, that integer must map to the same order or the request is rejected with 409 Conflict. This prevents a client with stale state from acting on the wrong order.

Two resource types are **not** in the mapping table. Products are addressed by SKU throughout. Channel integers are parsed directly from `channels.external_reference` when that column holds a numeric string, in `stockconnect-ce-channel.mapper.js` — so a channel whose external reference is non-numeric will not expose a `ChannelId` to CE clients.

13.8 Idempotency

The two surfaces differ, and this is the most consequential behavioural difference between them.

Surface	Header	Behaviour
<code>/api/v2</code> mutations	<code>Idempotency-Key</code> required	Missing header returns 400
<code>/api/v2/ce</code> routes delegating to core commands	Optional	A deterministic key is derived when absent
<code>/api/v2/ce</code> catalog and delivery-state routes	Optional	A key is computed but not passed to the ledger

Derivation when the header is absent:

```
export function resolveStockConnectCeIdempotencyKey(input) {
  const rawHeader = Array.isArray(input.header) ? input.header[0] : input.header;
  if (typeof rawHeader === 'string' && rawHeader.trim().length > 0) {
    return rawHeader.trim();
  }
  const digest = createHash('sha256')
    .update(`${input.tenantId}\0${input.routeId}\0${input.fingerprint}`)
    .digest('hex');
  return `ce-compat:${digest}`;
}
```

The fingerprint is a stable SHA-256 over the request body with sorted keys, so two byte-identical requests produce the same key and the second is recognised as a replay. This gives StockConnect at-least-once retry safety without requiring it to generate keys.

One subtlety worth recording: when a CE route delegates to a core command, the **ledger row** is written with the canonical `/api/v2/...` route identifier, not the `/api/v2/ce/...` path. The CE route identifier participates only in the hash that derives the key. This is intentional and documented in `docs/architecture/compatibility-idempotency.md`, and it means the same logical operation is deduplicated consistently regardless of which surface issued it.

The gap. `StockConnectCeCatalogCommand` and `StockConnectCeShipmentDeliveryCommand` receive a computed key but never call the idempotency service. A retried `POST /api/v2/ce/products` Or `PUT /api/v2/ce/offer/stock` can therefore repeat its side effects. For stock and price writes, which are absolute-value operations, a repeat is harmless. For `pushProducts`, which creates a product when the SKU is unknown, a genuine duplicate is prevented only by the tenant-unique SKU constraint.


13.9 Error mapping

Core errors are translated to the ChannelEngine envelope by `mapCoreErrorToExternalApiResponse`:

```
return {
  statusCode: mapped.statusCode,
  body: {
    Success: false,
    StatusCode: mapped.statusCode,
    Message: mapped.body.message,
    ...(validationErrors === undefined ? {} : { ValidationErrors: validationErrors }),
  },
};
```

Normalisation happens first: Zod issues and Fastify validation errors become `ValidationError`, a JSON parse failure becomes `MalformedRequestError`, and anything unrecognised becomes `InternalError`. The HTTP status comes from the `AppError`, so the status codes a CE client sees line up with the taxonomy in the API chapter. An integration test asserts that internal details are redacted rather than surfaced in `Message`.

13.10 Product data flow



Two catalog implementations exist and they are **not** the same code path.


Aspect	<code>/api/v2/products</code>	<code>/api/v2/ce/products</code>
Command	<code>CatalogCompatibilityCommand</code>	<code>StockConnectCeCatalogCommand</code>
Request schema	Typed CE product schema	<code>z.array(z.record(z.unknown()))</code>
Behaviour	Full upsert with content mapping and channel resolution	Ensures a product row exists by SKU only
Idempotency ledger	Used	Not used
Authorization checks	Explicit	Lighter
Response	Partial-success envelope from <code>buildCeMutationEnvelope</code>	CE envelope with per-item outcomes

An engineer debugging a "product pushed but fields missing" report on the CE surface should look here first: the CE push does not map rich product content.

Reading products also differs. `GET /api/v2/products` may require a content row to exist, whereas `GET /api/v2/ce/products` uses `mapProductToCeCatalogItem` and returns the product without requiring localised content.

13.11 Inventory data flow

Stock arriving from StockConnect is an **absolute value**, not a delta. The CE contract's `PUT /offer/stock` semantics are "set the sellable quantity to this number".




The stock location is resolved from the channel's `defaultStockLocationId`, falling back to a UUID stored in `configurationReference`. When neither is set the sync job is skipped with a recorded reason rather than failing — a common cause of "inventory never reaches the marketplace" reports.

Note the direction change at the bottom of the diagram. Stock written *in* through CE is stored against a stock location; stock sent *out* to a marketplace is read back as `available` at the channel's location. The two are connected through the database, not passed directly, which is why an inventory write and the resulting marketplace update are separated in time.

13.12 Pricing data flow

`PUT /api/v2/ce/offer` carries a price for a merchant product number. The CE catalog command defaults the currency to **SAR** when the payload omits it — a hard-coded default appropriate to the deployment region, worth knowing before assuming currency is always explicit. The price is written through the pricing contract as minor units with a validity window, then `price.created` or `price.updated` is recorded on the outbox, which plans a price sync job. The worker resolves the effective price through `getEffectivePrice` and the adapter converts minor units to whatever decimal format the marketplace expects.

13.13 Order lifecycle




Two behaviours in this machine regularly surprise newcomers and both are verified in code.

SHIPPED can be reached before anything physically ships.

`DefaultOrderFulfillmentService.evaluateOrderShipmentState` advances the order through **PROCESSING** and **READY_TO_SHIP** to **SHIPPED** once every non-cancelled line is fully allocated to shipments, in a single transaction. The order status reflects *allocation*, and the shipment's own status reflects physical dispatch.

DELIVERED and **RETURNED** are never set. Both transitions exist in `LEGAL_TRANSITIONS`, but no application code calls `Order.transitionTo(OrderStatus.DELIVERED)` or `RETURNED`. Delivering a shipment updates the shipment only; receiving a return updates the return and the order line's returned quantity, not the order status. Any consumer treating order status as a delivery signal will wait indefinitely — they should read shipment status instead.

The CE order payload maps Nexora status onto the ChannelEngine vocabulary, for example `CONFIRMED` becoming `IN_PROGRESS` and `DELIVERED` becoming `CLOSED`.



13.14 Delivery-state reporting

`PUT /api/v2/ce/shipments/:merchantShipmentNo/delivery-state` accepts a CE delivery status. Only `DELIVERED` and `CLOSED` are mapped to the internal deliver flow. Any other value returns a successful HTTP response carrying `Applied: false` and `Reason: 'STATUS_NOT_MAPPED'`. The call succeeded; nothing changed. Clients that only check the HTTP status will believe a state was recorded when it was not.

13.15 Invoice retrieval

`GET /api/v2/ce/orders/:merchantOrderNo/invoice` returns a PDF. If an `OrderInvoiceDocumentPort` implementation is wired, that provides the document. Otherwise `generateStockConnectCeInvoicePdf` produces one, and the file header states its scope plainly:

Nexora does not yet store marketplace tax invoices; this generator produces a readable PDF so StockConnect can parse invoice metadata.

The generated document is a single-page PDF containing an invoice number of the form `CE-<merchantOrderNo>`, an invoice date, the merchant order number and, when available, the channel order number. It contains no line items, no tax breakdown and no totals. It satisfies a parser looking for invoice metadata; it is not a tax document.

13.16 Outbound webhooks to StockConnect

Nexora can push order events to StockConnect rather than requiring polling. The mechanism reuses the general tenant webhook subsystem with a marker.

A webhook subscription whose `description` is exactly `stockconnect-ce-bridge` is recognised by `isStockConnectCeBridgeSubscription`. For such subscriptions the worker substitutes a CE-shaped payload strategy, and `buildStockConnectCeWebhookBody` renders the body as `{ "Content": [ ceOrder ] }` using the same order mapper as the polling endpoint — so a webhook payload and a poll response describe an order identically.

The bridged event types are exactly four:

```
order.created , order.confirmed , order.status_changed , order.cancelled
```

Delivery uses the standard signing and retry machinery described in the webhooks chapter, including the `x-Nexora-Signature` HMAC header. The description string is a magic value; changing it on a subscription silently reverts that subscription to the generic payload shape.

13.17 Readiness for the compatibility surface

Two readiness probes are registered on the API process, `external_compat_ce_routes` and its legacy alias `stockconnect_ce_compat`. Both run the same synchronous check, `assertExternalCompatibilityCeRoutesWired`, confirming that four dependencies are present on the route dependency object:

```
stockConnectCeOrderCompatibilityQuery , stockConnectCeCatalogCommand , stockConnectCeOrderInvoiceQuery and stockConnectCeChannelCompatibilityQuery .
```

What this tells you is that composition did not silently drop a CE handler — a real risk in a manually wired application. What it does **not** tell you is whether PostgreSQL is reachable, whether StockConnect is up, or whether any CE request has ever succeeded. Those are covered by the `postgres`, `redis` and `queue` probes and by actual traffic.

13.18 Operational workflows

Onboarding a StockConnect tenant. Create the tenant, create a channel and set its `defaultStockLocationId`, set the channel's `externalReference` to a numeric string if CE channel identifiers are required, create an API key with the scopes the integration needs, and give StockConnect the base URL plus the key. Optionally create a webhook subscription with description `stockconnect-ce-bridge` for push delivery.

Publishing catalogue changes. StockConnect pushes products, then stock and price. Each write records an integration event; the outbox publisher enqueues sync jobs; the worker calls the marketplace adapter. Confirm with `GET /api/v2/ce/channels/:channelId/products`, which reports `PUBLISHED`, `DISABLED` or `NOTPUBLISHED` derived from the offer status.

Processing orders. StockConnect polls `GET /api/v2/ce/orders`, acknowledges, creates shipments, and reports delivery state. Note that the CE poll endpoint does **not** filter to new orders — see the limitations below.

Recovering from a failed sync. Sync failures do not fail the originating API call. Inspect worker logs for the catalog sync job, check whether the offer has an `externalReference`, whether the channel has a stock location, and whether the marketplace connection tests successfully via `POST /api/v1/channels/:channelId/marketplace-connection/test`.

13.19 Troubleshooting

Symptom	Likely cause	Where to look	Action
401 on every CE call	Query key used on <code>/api/v2/*</code> rather than <code>/api/v2/ce/*</code>	<code>authentication.plugin.js</code> , <code>merchant-compat-route-prefix.js</code>	Move to the <code>/ce</code> prefix or send <code>x-api-key</code>
401 intermittently	Both <code>Authorization</code> and an API key sent	Auth plugin dual-credential check	Send exactly one credential
400 <code>VALIDATION_FAILED</code> on a v2 mutation	Missing <code>Idempotency-Key</code>	<code>requireIdempotencyKey</code>	Add the header; CE routes derive one automatically
409 <code>CONFLICT</code> on an order operation	<code>OrderId</code> integer does not match <code>MerchantOrderNo</code>	<code>compatibility-external-id-resolution.js</code>	Re-fetch the order; do not reuse cached identifiers
409 <code>IDEMPOTENCY_CONFLICT</code>	Same key, different body	<code>idempotency_records</code>	Use a new key for a genuinely different request
429	Compatibility rate policy exceeded	<code>rate-limit-policies.js</code>	Honour <code>Retry-After</code> ; reads 120/min, mutations 60/min per credential
Product created but fields missing	CE <code>pushProducts</code> only ensures the SKU exists	<code>stockconnect-ce-catalog-command.js</code>	Use <code>/api/v2/products</code> for full upsert, or extend the CE command
Stock accepted but marketplace unchanged	Offer has no <code>externalReference</code> , or channel has no stock location	<code>execute-catalog-sync-job.js</code> , <code>resolve-channel-stock-location-id.js</code>	Set the offer external reference and the channel default stock location
Delivery state accepted but nothing changed	Status other than <code>DELIVERED</code> or <code>CLOSED</code>	<code>stockconnect-ce-shipment-delivery-command.js</code>	Check <code>Applied</code> and <code>Reason</code> in the response body
Order never reaches <code>DELIVERED</code>	No code drives that transition	<code>order-status.js</code>	Read shipment status instead

Symptom	Likely cause	Where to look	Action
CE poll returns old orders	Poll has no <code>NEW</code> -only filter	<code>stockconnect-ce-order-compatibility-query.js</code>	Filter client-side, or use <code>/api/v2/orders/new</code>
<code>ChannelId</code> absent from CE payloads	Channel <code>externalReference</code> is not numeric	<code>stockconnect-ce-channel.mapper.js</code>	Set a numeric external reference
Duplicate side effects after retry	CE catalog routes bypass the idempotency ledger	<code>stockconnect-ce-catalog-command.js</code>	Treat these as at-least-once; prefer absolute-value operations
Readiness green but CE failing	The probe only checks wiring	<code>stockconnect-ce-readiness.js</code>	Check <code>postgres</code> , <code>redis</code> , <code>queue</code> probes and logs

13.20 Implementation status

Implemented and covered by automated tests

- The complete `/api/v2` Merchant surface with required idempotency
- The complete `/api/v2/ce` route set, asserted by a dedicated route-inventory integration test
- CE query-parameter and `X-CE-KEY` credential resolution, restricted to the `/ce` prefix
- Delegation of orders, shipments, cancellations and returns to shared compatibility commands
- `compat_v2` integer identifier allocation, enrichment and strict inbound resolution
- CE order polling, acknowledgement, channel listing, channel product status and product listing
- The CE error envelope with internal-detail redaction
- The `stockconnect-ce-bridge` webhook payload builder
- CE readiness probe wiring

Partially implemented

Area	Limitation
Invoice PDF	Metadata-only placeholder; no line items, tax or totals
CE catalog push	Ensures a product row exists by SKU; does not map rich content like <code>/api/v2</code> does
CE idempotency	Catalog and delivery-state routes compute a key but never persist it
CE order poll	No default <code>NEW</code> -only filter, unlike <code>/api/v2/orders/new</code>
Channel identifiers	Only exposed when <code>external_reference</code> is numeric
Delivery state	Only <code>DELIVERED</code> and <code>CLOSED</code> are mapped
Return receive	Partial accept or reject across lines is not supported
CE field coverage	Several accepted fields — tracking extras, cancellation reasons, refund amounts — are validated but ignored

Not implemented

- Any outbound call to ChannelEngine; there is no client, no credentials, no configuration
- ChannelEngine asynchronous job, batch and queue administration APIs, recorded as out of scope
- Channel order ingest on the CE prefix; ingest exists only at `POST /api/v2/orders`
- Merchant-only routes on the CE prefix, listed earlier in this chapter
- Product and channel entries in the integer identifier table

Verification status. Automated unit and integration coverage for the compatibility surface is extensive — twenty-two unit test files and more than twenty-eight integration test files, including eight specifically named for StockConnect CE. However, `docs/architecture/stockconnect-ce-production-cutover.md` states explicitly that **live StockConnect staging and production traffic has not been exercised**, marking those checks as blocked by environment. No claim of production verification should be made on the strength of the in-repository tests alone.

13.21 Future considerations


These are recommendations, not implemented behaviour.

Area	Suggested direction
CE catalog idempotency	Route <code>StockConnectCeCatalogCommand</code> through the same idempotency service the core commands use
CE product push parity	Reuse <code>CatalogCompatibilityCommand</code> 's content mapping so both surfaces behave identically
Invoice documents	Implement <code>OrderInvoiceDocumentPort</code> against real invoice data or a marketplace provider
Order poll filtering	Add an explicit status filter to the CE poll so clients need not filter locally
Order delivery status	Decide whether shipment delivery should drive the order aggregate, and implement or document the decision
Webhook bridge marker	Replace the magic <code>description</code> string with an explicit subscription attribute
Live verification	Complete the staging exercise the cutover runbook identifies as outstanding

14 Product, inventory, pricing and order workflows


This chapter follows four end-to-end flows through the layers, to show how the modules described earlier cooperate.

14.1 Publishing a product to a marketplace



The order of these calls matters. Activating an offer configured to resolve pricing will fail if no effective price exists for the channel and currency, which is a deliberate guard against publishing an unpriced listing.

14.2 Inventory movement through an order's life



At every step the invariant `available = on_hand - reserved` holds, and every step leaves an immutable movement row. Reconstructing why a balance is what it is is therefore always possible from `inventory_movements`.

14.3 Channel-fulfilled orders

When a marketplace has already shipped the goods, reserving stock would be wrong — the stock has left. `POST /api/v2/orders/channel-fulfilled` creates the order with an inventory consumption mode of `record_sale_only`. No reservation is taken at creation, and when the associated shipment is marked shipped, `fulfillInventoryWhenShipmentShipped` calls `recordSale` instead of `fulfillReservedForShipment`. The net effect on `on_hand` is the same; the difference is that `available` drops at shipment rather than at order creation.

14.4 Request equivalence in practice

Consider a marketplace webhook delivered twice. The second delivery calls `CreateChannelOrder` with the same channel and the same `externalOrderReference`. Rather than creating a duplicate or failing, the use case loads the existing order and runs `isEquivalentChannelOrderRequest`, comparing channel, currency, discount, tax and shipping totals, a sorted line snapshot and a normalised customer record.

Outcome	Response
Payloads equivalent	The existing order is returned with <code>ingestionOutcome: 'duplicate'</code>
Payloads differ	<code>ConflictError</code> , because the same external reference now describes different data

The same pattern guards shipments, returns and cancellations on their `externalReference`. It complements, rather than replaces, `Idempotency-Key`: idempotency keys protect identical retries of the same HTTP call, while equivalence protects against duplicate ingestion of the same business event arriving by different routes.

15 Queue and background processing


15.1 Queues and jobs

Queue	Job name	Handler behaviour
integration-events	publish-integration-event	Routes an outbox event to its in-process consumers
webhook-deliveries	deliver-webhook	Signs and POSTs a webhook, classifies the response
channel-catalog-sync	run-catalog-sync	Executes one product, offer, price or inventory sync
marketplace-order-lifecycle	process-marketplace-lifecycle	Processes webhook events, inbound lifecycle commands and outbound operations

All four are BullMQ queues backed by Redis, sharing one configuration:

Setting	Environment variable	Default
Attempts	QUEUE_DEFAULT_ATTEMPTS	5
Backoff base, exponential	QUEUE_BACKOFF_BASE_MS	1000
Worker concurrency	QUEUE_WORKER_CONCURRENCY	5
Job lock duration	QUEUE_JOB_TIMEOUT_MS	30000
Key prefix	QUEUE_PREFIX	nexora-queue
Redis URL	QUEUE_REDIS_URL , falling back to REDIS_URL	—

15.2 Worker architecture



Each job executes inside an `AsyncLocalStorage` request context carrying a correlation identifier, so worker log lines can be joined to the API request that originally produced the event.

15.3 Retry semantics

Three distinct retry mechanisms coexist, and confusing them makes failures hard to diagnose.

Mechanism	Used by	Behaviour
BullMQ attempts with exponential backoff	All queues	Generic failure retry up to <code>QUEUE_DEFAULT_ATTEMPTS</code>

Mechanism	Used by	Behaviour
<code>moveToDelayed</code> with a computed delay	Webhook delivery, catalog sync	The handler decides the delay — for example honouring <code>Retry-After</code>
Permanent failure, no retry	Catalog sync, webhook delivery	Adapter classified the error as permanent; logged and dropped

A catalog sync job raising `CatalogSyncRetryError` is delayed and retried; an adapter authentication or validation error is classified permanent and logged without a retry, because retrying bad credentials or an invalid payload cannot succeed.

15.4 Schedulers

Scheduler	Process	Trigger	Purpose
Outbox publisher	API	<code>OUTBOX_POLL_INTERVAL_MS</code> , default 1000	Move committed events onto the queue
Retention cleanup	Worker	<code>RETENTION_CLEANUP_INTERVAL_MS</code> , default 3600000	Delete aged outbox, inbox, idempotency and delivery rows
Catalog sync reconciliation	Worker	<code>CATALOG_SYNC_RECONCILIATION_INTERVAL_MS</code> , default 86400000, only when enabled	Re-plan sync jobs for offers that may have drifted

Reconciliation is **disabled by default** (`CATALOG_SYNC_RECONCILIATION_ENABLED` defaults to false). It exists because event-driven sync can miss changes when a job is permanently failed or an event is lost; the periodic sweep re-plans jobs for a bounded batch of offers per tick.

16 Marketplace integrations

16.1 Adapter ports

Provider code is reached only through five interfaces, so adding a marketplace means implementing ports rather than modifying core code.

Port	Responsibility	Key methods
MarketplaceCatalogAdapter	Outbound catalogue	<code>testConnection</code> , <code>syncProduct</code> , <code>syncOffer</code> , <code>syncPrice</code> , <code>syncInventory</code> , <code>getCapabilities</code>
MarketplaceAdapterRuntime	Credentials and configuration for one channel	Produced by <code>createForSync</code>
MarketplaceOrderAdapter	Inbound orders	<code>listOrders</code> , <code>fetchOrder</code> , <code>normalizeWebhookOrder</code> , <code>normalizeLifecycleCommand</code>
MarketplaceWebhookAdapter	Inbound webhooks	<code>authenticateWebhookRequest</code> , <code>normalizeWebhookEvent</code>
MarketplaceOutboundOrderLifecycleAdapter	Outbound order actions	<code>cancelOrder</code> , <code>refundOrder</code> , <code>createFulfillment</code>

16.2 Support matrix

Read this table before promising a capability to a stakeholder. It reflects adapter code, not marketing.

Capability	Amazon	Noon	Namshi	Shopify
Connection test	Yes	Yes	Yes	Yes
Product sync	Partial — SKU mapping and availability, no full listing creation	No — throws unsupported	No — throws unsupported	Yes
Offer sync	Yes	Yes	Yes	Yes
Price sync	Yes	Yes	Yes	Yes
Inventory sync	Yes	Yes	Yes	Yes
Order polling	No	No	No	Yes
Order ingestion via lifecycle or webhook	Partial	Partial	Partial	Yes
Inbound lifecycle events	Yes	Yes	Yes	Partial
Outbound cancel	No	No	No	Yes
Outbound refund	No	No	No	Yes
Outbound fulfilment	Yes	Yes	Yes	Yes

Capability	Amazon	Noon	Namshi	Shopify
Outbound returns	No	No	No	No
Inbound webhooks	Partial — SNS handshake unsupported	Yes	Yes	Yes when registered


Noon and Namshi require SKUs to already exist on the marketplace; their `syncProduct` throws `MarketplaceUnsupportedError` by design, with a comment stating the precondition. Shopify is the only provider with a complete inbound and outbound order lifecycle.

16.3 Authentication per provider

Provider	Mechanism
Amazon SP-API	LWA refresh token exchanged for an access token, requests signed AWS4-HMAC-SHA256 for <code>execute-api</code>
Noon	RS256 JWT login producing a cookie session
Namshi	RS256 JWT login producing a cookie session
Shopify	Admin API access token; webhooks verified by base64 HMAC-SHA256

Credentials are stored encrypted in `marketplace_connections.credentials_ciphertext` and decrypted by the runtime factory when a sync job runs. They are never returned by the read endpoint, and connection-test errors pass through `sanitizeMarketplaceConnectionTestError`, which strips secret-looking substrings and truncates to 500 characters.


16.4 Catalog synchronisation pipeline



The job identifier is the coalescing mechanism. Inventory jobs use a key of tenant, channel, target and entity, so a burst of stock changes for one product collapses into a single pending job rather than a queue of redundant updates. Price job keys additionally include currency.

The rate limiter allows 120 sync operations per minute per tenant-and-channel pair. This limit is **hard-coded** in `rate-limit-policies.js` and has no environment override, unlike the reconciliation settings.


16.5 Order ingestion



Only three normalized statuses are ingestible as new orders: `pending`, `confirmed` and `unknown`. An order arriving already cancelled or already fulfilled is not created through this path.

Lifecycle application is deduplicated through the idempotency service using the marketplace's `externalEventId` as the key under route identifier `marketplace-order-lifecycle.apply`, so a marketplace redelivering the same event does not apply it twice.

16.6 Inbound webhook flow



The route is public because the tenant is unknown until the ingress token is resolved. The token is stored only as a SHA-256 hash, and the lookup uses a `SECURITY DEFINER` function so it can run before any tenant context exists. Provider authentication then applies on top: Shopify verifies an HMAC over the raw body, Noon and Amazon parse and validate the envelope, and Namshi relies on the ingress token alone.

The raw body must be preserved byte-for-byte for HMAC verification, which is why the route captures `marketplaceWebhookRawBody` rather than working from the parsed JSON.

17 Webhooks

This chapter covers **outbound** webhooks — Nexora notifying a tenant's endpoint. Inbound marketplace webhooks are covered in the previous chapter.

17.1 Subscription model


A subscription holds a destination URL, an encrypted secret, and a list of event types. Event types are validated at creation against the catalogue's external-delivery allowlist, so a subscription cannot be created for an event that will never be dispatched. The destination URL is validated against the SSRF rules described below.

17.2 Event catalogue

Twenty-eight integration event types exist in `src/shared/events/event-catalog.js` :

Domain	Events
Orders	<code>order.created</code> , <code>order.confirmed</code> , <code>order.status_changed</code> , <code>order.cancelled</code>
Shipments	<code>shipment.created</code> , <code>shipment.shipped</code> , <code>shipment.delivered</code> , <code>shipment.cancelled</code> , <code>shipment.status_changed</code>
Cancellations	<code>cancellation.created</code> , <code>cancellation.completed</code>
Returns	<code>return.created</code> , <code>return.status_changed</code>
Products	<code>product.created</code> , <code>product.updated</code> , <code>product.status_changed</code>
Inventory	<code>inventory.inventory_changed</code> , <code>inventory.inventory_reserved</code> , <code>inventory.inventory_released</code>
Offers	<code>offer.created</code> , <code>offer.updated</code> , <code>offer.status_changed</code>
Channels	<code>channel.created</code> , <code>channel.updated</code> , <code>channel.status_changed</code>
Prices	<code>price.created</code> , <code>price.updated</code> , <code>price.changed</code>

17.3 Delivery



Headers sent with every delivery:

Header	Content
X-Nexora-Signature	v1=<hex> HMAC-SHA256 of the request body using the subscription secret
X-Nexora-Event	Event type
X-Nexora-Event-Id	Event identifier
X-Nexora-Delivery-Id	Delivery attempt identifier

Response classification is explicit: 408, 425, 429 and any 5xx are retryable; all other non-2xx responses are permanent failures. A `Retry-After` header on a retryable response is honoured up to `WEBHOOK_DELIVERY_MAX_RETRY_AFTER_SECONDS`. When attempts are exhausted the delivery is marked dead-lettered and remains queryable through `GET /api/v1/webhooks/:webhookId/deliveries`.

17.4 SSRF protection

`validateOutboundWebhookUrl` runs at both subscription time and delivery time. It requires HTTPS, rejects URLs containing embedded credentials, resolves all A and AAAA records, and blocks private, loopback, link-local and cloud metadata address ranges. Validating again at delivery time defends against DNS rebinding between registration and use.

Destinations are logged through `sanitizeWebhookDestinationForLog`, which records only host and pathname, so query-string tokens in a destination URL do not reach the logs.

17.5 Secret rotation

POST `/api/v1/webhooks/:webhookId/rotate-secret` requires the `webhooks.manage` permission plus a valid MFA step-up session, and is rate limited under the `webhook.secret-rotate` policy.

18 External services

Service	Direction	Purpose	Configuration by name
PostgreSQL	Outbound	Source of truth	<code>DATABASE_URL</code> , <code>DATABASE_MIGRATION_URL</code> , pool and timeout variables
Redis	Outbound	Cache, locks, rate limiting, queue	<code>REDIS_URL</code> , <code>REDIS_KEY_PREFIX</code> , <code>QUEUE_REDIS_URL</code>
S3-compatible storage	Outbound	Object storage provider, wired but unused by domain modules	<code>STORAGE_BUCKET</code> , <code>STORAGE_REGION</code> , <code>STORAGE_ENDPOINT</code> , <code>STORAGE_ACCESS_KEY_ID</code> , <code>STORAGE_SECRET_ACCESS_KEY</code> , <code>STORAGE_FORCE_PATH_STYLE</code>
OTLP collector	Outbound	Trace export	<code>TRACING_ENABLED</code> , <code>OTEL_EXPORTER_OTLP_ENDPOINT</code> , <code>OTEL_SERVICE_NAME</code>
Amazon SP-API	Outbound	Catalogue, inventory, price, fulfilment	Per-channel credentials; <code>AMAZON_LWA_TOKEN_URL</code> optional
Noon	Outbound	Catalogue, inventory, price, fulfilment	Per-channel credentials
Namshi	Outbound	Catalogue, inventory, price, fulfilment	Per-channel credentials
Shopify	Both	Full catalogue and order lifecycle	Per-channel credentials; <code>SHOPIFY_ADMIN_API_VERSION</code> optional
StockConnect	Inbound plus outbound webhooks	Consumes the CE compatibility API	Authenticated by tenant API key; no Nexora-side configuration
Tenant webhook endpoints	Outbound	Event notification	Per-subscription URL and secret

18.1 Outbound HTTP client

All outbound HTTP goes through `src/infrastructure/http/fetch-http-client.js` , which applies a default timeout of `HTTP_CLIENT_TIMEOUT_MS` , propagates the `x-request-id` header, and constructs a W3C `traceparent` header from the active OpenTelemetry trace. `GET` and `HEAD` are retried once on timeout; `POST` is not retried automatically, because the client cannot know whether the request was idempotent. URLs are redacted before logging.

18.2 A configuration inconsistency worth knowing

`create-application.js` reads `config.marketplace.noonApiBaseUrl` and `config.marketplace.noonUserAgent` , but `src/app/config/config.js` only populates `shopifyAdminApiVersion` and `amazonLwaTokenUrl` , and no `NOON_*` variables exist in the schema. Those two values are therefore always `undefined` at runtime, and the Noon

adapters fall back to the in-code defaults in `noon-config.js`. This is harmless today but will confuse anyone who tries to override the Noon base URL by setting an environment variable.

19 Error handling and observability

19.1 Error normalisation

Every error passing through the API converges on one path. `src/app/errors/error-mapper.js` normalises the input: an `AppError` passes through, a `ZodError` becomes a `ValidationError` with field issues, a `SyntaxError` becomes a `MalformedRequestError`, a Fastify `FST_ERR_VALIDATION` becomes a `ValidationError`, and anything else becomes an `InternalServerError`. The envelope builder then produces the JSON described in the API chapter.

Operational errors — those the `AppError` marks as expected — are logged at `warn`; everything else is logged at `error`. This keeps alerting signal high, because a client sending an invalid payload does not raise the same log level as an unexpected exception.

19.2 Structured logging

Pino emits JSON with base fields `service` and `env`, enriched per request from `AsyncLocalStorage` with `requestId`, `traceId`, `tenantId` and `userId`. Redaction is configured at two levels: `PINO_REDACT_PATHS` covers header and body paths such as `authorization`, `cookies`, `x-api-key`, `x-ce-key`, tokens and passwords, and `redactDeep` handles arbitrary nested payloads, masking PII such as email addresses to the form `j***@domain`.

Access logs are written on `onResponse` with method, route, status code and elapsed time.

19.3 Metrics

Exposed at `/internal/metrics` on the API process and on the worker's observability port.

Group	Metric names
HTTP	<code>http_requests_total</code> , <code>http_request_duration_seconds</code> , <code>http_request_errors_total</code>
Database	<code>db_pool_connections</code> , <code>db_query_duration_seconds</code>
Redis	<code>redis_operations_total</code> , <code>redis_operation_duration_seconds</code>
Queue	<code>queue_jobs_total</code> , <code>queue_job_duration_seconds</code>
Security	<code>rate_limit_hits_total</code> , <code>auth_events_total</code>
Commerce	<code>commerce_operations_total</code>
Webhooks	<code>webhook_delivery_attempts_total</code> , <code>webhook_delivery_duration_seconds</code>
Retention	<code>nexora_retention_cleanup_deleted_total</code> , <code>_runs_total</code> , <code>_failures_total</code> , <code>_duration_seconds</code>
Catalog sync	<code>channel_catalog_sync_jobs_total</code> , <code>channel_catalog_reconciliation_*</code>
Marketplace	<code>marketplace_order_ingestion_total</code> , <code>marketplace_webhook_total</code> , <code>marketplace_order_lifecycle_total</code> , <code>marketplace_lifecycle_worker_jobs_total</code>

Node default metrics are also collected. The database pool gauge is refreshed on a five-second interval in both processes.

19.4 Tracing


OpenTelemetry is enabled when `TRACING_ENABLED` is true and `OTEL_EXPORTER_OTLP_ENDPOINT` is set — configuration validation rejects the first without the second. HTTP, `pg` and `ioredis` are instrumented, and the HTTP instrumentation ignores `/health` and `/internal/metrics` to avoid flooding traces with probe traffic. Trace identifiers flow into log lines and into outbound `traceparent` headers.

19.5 Health and readiness

Endpoint	Process	Semantics
GET <code>/health/live</code>	API	Always <code>{ status: 'ok' }</code> ; does not reflect shutdown
GET <code>/health/ready</code>	API	200 with checks, or 503 when any probe fails or shutdown started
GET <code>/health/live</code>	Worker	503 <code>{ status: 'shutting_down' }</code> during drain
GET <code>/health/ready</code>	Worker	Worker probe set

Use readiness, not liveness, for load-balancer removal on the API process.

19.6 Diagnostic workflow



The `requestId` returned in every error envelope and in the `x-request-id` header is the primary correlation key; for background work, the correlation identifier carried into the job context serves the same purpose.

20 Development environment

20.1 Prerequisites

Requirement	Version
Node.js	24 or later, per engines
npm	10.9.2, per packageManager
Docker and Docker Compose	For PostgreSQL, Redis and MinIO

20.2 First run

```
npm install
cp .env.example .env           # then fill in the required values
docker compose up -d           # postgres, redis, minio
npm run migrate                 # apply migrations
npm run dev                    # API on SERVER_PORT, default 3000
npm run dev:worker              # worker, in a second terminal
```

The Compose stack publishes PostgreSQL on host port **5433**, mapped to 5432 in the container. Note that docs/operations/README.md states 5432; the Compose file is authoritative.

20.3 Required configuration

Only a few variables have no usable default:

Variable	Why it is required
DATABASE_URL	No default connection string
REDIS_URL	No default
STORAGE_BUCKET , STORAGE_ACCESS_KEY_ID , STORAGE_SECRET_ACCESS_KEY	Storage provider is constructed at bootstrap
AUTH_JWT_SECRET , or both AUTH_JWT_PRIVATE_KEY and AUTH_JWT_PUBLIC_KEY	Enforced by cross-field validation

Configuration is validated at startup and the process exits on a violation. Cross-field rules enforced in assertConsistency :

- DATABASE_POOL_MIN must not exceed DATABASE_POOL_MAX
- TRACING_ENABLED requires OTEL_EXPORTER_OTLP_ENDPOINT
- In production, SERVER_CORS_ORIGINS must not contain *
- A JWT secret or a complete key pair must be present
- AUTH_MFA_ENCRYPTION_KEY must be exactly 32 characters

- `DATABASE_MIGRATION_URL` must parse as a URL when set
- `WEBHOOK_DELIVERY_TIMEOUT_MS` must not exceed `WEBHOOK_DELIVERY_LEASE_SECONDS × 1000`

20.4 Configuration reference by concern

Names only; never commit values.

Concern	Variables
Runtime	<code>NODE_ENV</code> , <code>APP_NAME</code> , <code>APP_INSTANCE_NAMESPACE</code>
HTTP server	<code>SERVER_HOST</code> , <code>SERVER_PORT</code> , <code>SERVER_SHUTDOWN_TIMEOUT_MS</code> , <code>SERVER_BODY_LIMIT_BYTES</code> , <code>SERVER_CORS_ORIGINS</code> , <code>SERVER_TRUST_INCOMING_REQUEST_ID</code> , <code>SERVER_TRUST_PROXY</code>
Database	<code>DATABASE_URL</code> , <code>DATABASE_MIGRATION_URL</code> , <code>DATABASE_POOL_MAX</code> , <code>DATABASE_POOL_MIN</code> , <code>DATABASE_CONNECTION_TIMEOUT_MS</code> , <code>DATABASE_IDLE_TIMEOUT_MS</code> , <code>DATABASE_STATEMENT_TIMEOUT_MS</code> , <code>DATABASE_QUERY_TIMEOUT_MS</code> , <code>DATABASE_SSL</code>
Redis	<code>REDIS_URL</code> , <code>REDIS_KEY_PREFIX</code> , <code>REDIS_CONNECT_TIMEOUT_MS</code> , <code>REDIS_COMMAND_TIMEOUT_MS</code>
Queue	<code>QUEUE_REDIS_URL</code> , <code>QUEUE_PREFIX</code> , <code>QUEUE_DEFAULT_ATTEMPTS</code> , <code>QUEUE_BACKOFF_BASE_MS</code> , <code>QUEUE_WORKER_CONCURRENCY</code> , <code>QUEUE_JOB_TIMEOUT_MS</code>
Storage	<code>STORAGE_ENDPOINT</code> , <code>STORAGE_REGION</code> , <code>STORAGE_BUCKET</code> , <code>STORAGE_ACCESS_KEY_ID</code> , <code>STORAGE_SECRET_ACCESS_KEY</code> , <code>STORAGE_FORCE_PATH_STYLE</code>
Outbox and idempotency	<code>OUTBOX_BATCH_SIZE</code> , <code>OUTBOX_POLL_INTERVAL_MS</code> , <code>OUTBOX_MAX_ATTEMPTS</code> , <code>IDEMPOTENCY_TTL_SECONDS</code>
Retention	<code>OUTBOX_RETENTION_DAYS</code> , <code>INBOX_RETENTION_DAYS</code> , <code>IDEMPOTENCY_RETENTION_DAYS</code> , <code>WEBHOOK_DELIVERY_RETENTION_DAYS</code> , <code>RETENTION_CLEANUP_BATCH_SIZE</code> , <code>RETENTION_CLEANUP_INTERVAL_MS</code>
Catalog reconciliation	<code>CATALOG_SYNC_RECONCILIATION_ENABLED</code> , <code>_INTERVAL_MS</code> , <code>_OFFER_BATCH_SIZE</code> , <code>_MAX_JOBS_PER_TICK</code>
Observability	<code>LOG_LEVEL</code> , <code>LOG_PRETTY</code> , <code>METRICS_ENABLED</code> , <code>TRACING_ENABLED</code> , <code>OTEL_EXPORTER_OTLP_ENDPOINT</code> , <code>OTEL_SERVICE_NAME</code>
Worker observability	<code>WORKER_OBSERVABILITY_HTTP_ENABLED</code> , <code>WORKER_OBSERVABILITY_HOST</code> , <code>WORKER_OBSERVABILITY_PORT</code> , <code>WORKER_OTEL_SERVICE_NAME</code>
Outbound HTTP and webhooks	<code>HTTP_CLIENT_TIMEOUT_MS</code> , <code>WEBHOOK_DELIVERY_TIMEOUT_MS</code> , <code>WEBHOOK_DELIVERY_LEASE_SECONDS</code> , <code>WEBHOOK_DELIVERY_MAX_RETRY_AFTER_SECONDS</code>
Documentation	<code>DOCS_ENABLED</code>
Authentication	<code>AUTH_JWT_SECRET</code> , <code>AUTH_JWT_PRIVATE_KEY</code> , <code>AUTH_JWT_PUBLIC_KEY</code> , <code>AUTH_ACCESS_TOKEN_TTL_SECONDS</code> , <code>AUTH_REFRESH_TOKEN_TTL_SECONDS</code> , <code>AUTH_MFA_ENCRYPTION_KEY</code> , <code>AUTH_STEP_UP_TTL_SECONDS</code> , <code>AUTH_PASSWORD_MIN_LENGTH</code> , <code>AUTH_PASSWORD_MAX_LENGTH</code>
Marketplace adapters	<code>SHOPIFY_ADMIN_API_VERSION</code> , <code>AMAZON_LWA_TOKEN_URL</code>

There are **no** StockConnect or ChannelEngine environment variables. StockConnect authenticates with an ordinary tenant API key.

20.5 npm scripts

Script	Action
<code>dev / dev:worker</code>	Run the API or worker with <code>--watch</code>
<code>start / start:worker</code>	Run without watch
<code>migrate / migrate:status</code>	Apply or report migrations
<code>backfill:external-ids</code>	Populate compatibility integer identifiers for existing rows
<code>build</code>	Syntax check every file with <code>node --check</code>
<code>lint / lint:fix</code>	ESLint
<code>format / format:check</code>	Prettier
<code>arch:check</code>	dependency-cruiser boundary enforcement
<code>test / test:unit</code>	Unit suite
<code>test:integration</code>	Integration suite
<code>test:all</code>	Both
<code>verify</code>	<code>lint</code> then <code>arch:check</code> then <code>test:unit</code>
<code>docs:route-inventory</code>	Regenerate the route inventory
<code>docs:openapi-coverage</code>	Dump the OpenAPI snapshot and verify route coverage

20.6 Common local problems

Symptom	Cause	Fix
Process exits immediately with a configuration error	A required variable is missing or a cross-field rule failed	Read the message; it names the failing rule
<code>/health/ready</code> returns 503 with <code>storage: failed</code>	MinIO is not running or the bucket does not exist	<code>docker compose up -d</code> , confirm <code>minio-init</code> completed
Login returns 429	The <code>auth.login</code> policy allows 10 attempts per 15 minutes	Wait, or clear the Redis key
Any rate-limited call fails when Redis is down	The limiter fails closed by design	Restore Redis
Connection refused on port 5432	Compose publishes 5433	Use 5433 in <code>DATABASE_URL</code>

Symptom	Cause	Fix
Jobs never execute	The worker process is not running	Start <code>npm run dev:worker</code>
<code>/docs</code> returns 404	<code>DOCS_ENABLED</code> is false	Enable it in the local environment

21 Testing strategy

21.1 Suites

Project	Location	Files	External dependencies
unit	tests/unit/	136	None
integration	tests/integration/	73	PostgreSQL, Redis, MinIO

Integration tests run serially with a 30-second test timeout and a 60-second hook timeout.

`tests/integration/setup.js` loads `.env` without overriding existing values and supplies local defaults.

`tests/integration/helpers.js` builds a shared infrastructure singleton, and most tests then call

`createApplication(infra)` and use Fastify's `inject()` rather than opening a socket.

21.2 Coverage by area

Area	Unit files	What is covered
Marketplaces	28	All four provider adapters, signing, token providers, error mapping, contracts
Compatibility	22	CE mappers, external identifier enrichment and resolution, idempotency derivation, invoice generation, error mapping
Channel catalog sync	16	Job planning, identifiers, reconciliation, rate limiting, per-target sync
Webhooks	11	Dispatch, signing, retry delay, URL validation, response classification
Marketplace order ingestion	10	Schemas, lifecycle service, normalization, job execution
Cross-cutting	13	Configuration, errors, validation, pagination, money, redaction, retention
Authorization	4	Permission patterns, effective permissions, service
Orders	4	Status machine, command service
Identity	3	Email, refresh rotation, auth failures
Workers	3	Job handlers, event router, observability HTTP
Architecture	2	Asserts dependency-cruiser rule names exist
Others	1–2 each	Tenants, shipments, returns, cancellations, events, products, channels, API keys, security

Integration coverage is strongest around compatibility, with more than twenty-eight files including eight named for StockConnect CE, and solid across authentication, tenancy, commerce, webhooks and marketplace flows.

21.3 Running tests

```
npm run test:unit          # fast, no dependencies
docker compose up -d      # required for the next command
npm run test:integration
npm run test:all
npm run verify             # lint + arch:check + unit
```

21.4 Verified results

At the time of writing, `npm run test:unit` completed with **623 tests passing across 136 files**. The integration suite was not executed in full for this document, so no claim is made about its current pass rate. One known environment-dependent failure exists: `tests/integration/http.test.js` asserts `/health/ready` returns 200, which fails when a local dependency such as MinIO is unavailable.

21.5 Gaps


Gap	Note
No CI workflow in the repository	<code>.github/workflows/</code> does not exist
<code>verify</code> excludes integration tests, Prettier and the syntax check	Intentional for speed; run them separately before release
No coverage thresholds configured	Coverage is not measured in the repository
Live StockConnect verification absent	Stated in the cutover runbook
Real marketplace APIs are not called in tests	Adapters are exercised against fixtures
Documentation drift	<code>tests/README.md</code> references TypeScript and a typecheck step that do not exist

22 Deployment and operations

22.1 Container image

The `Dockerfile` is a multi-stage build on Node 24 Alpine. The runtime stage installs production dependencies only, runs as the unprivileged `node` user, uses `tiny` as PID 1 for correct signal handling, exposes port 3000, and defaults to `node src/app/main.js`. The worker uses the same image with the command overridden to `node src/workers/main.js`.

22.2 Process topology



API replicas scale with request volume; worker replicas scale with job volume. Both are stateless. Scheduled work is safe to run on multiple worker replicas because the retention job takes a distributed lock, and catalog sync jobs are coalesced by deterministic job identifiers.

22.3 Probe configuration

Purpose	API	Worker
Liveness	<code>GET /health/live</code> on <code>SERVER_PORT</code>	<code>GET /health/live</code> on <code>WORKER_OBSERVABILITY_PORT</code>
Readiness	<code>GET /health/ready</code> on <code>SERVER_PORT</code>	<code>GET /health/ready</code> on the worker port
Metrics scrape	<code>GET /internal/metrics</code>	Same path on the worker port

The worker observability listener binds `127.0.0.1` by default. In-cluster probes require `WORKER_OBSERVABILITY_HOST=0.0.0.0`, and ADR-026 is explicit that this port must never be exposed through a public load balancer.

22.4 Migrations in production

Run migrations as a separate step before rolling out new application pods, using a privileged role. Set `DATABASE_MIGRATION_URL` to that role's connection string and `DATABASE_URL` to the restricted `nexora_app` role, so that runtime traffic is subject to row-level security while DDL is not.

Migrations are forward-only. There are no down migrations; a bad migration is corrected by writing a new one.

22.5 Shutdown behaviour

On `SIGTERM` the process marks itself not ready, then closes listeners, schedulers, workers, queue, Redis, database and tracing in order, bounded by `SERVER_SHUTDOWN_TIMEOUT_MS`. The orchestrator's termination grace period should exceed that value, otherwise the process is killed mid-drain.

22.6 Security hardening checklist for production

Derived from the verified gaps recorded in this document.

Item	Reason
Override <code>AUTH_MFA_ENCRYPTION_KEY</code>	The schema default is a known value
Change the <code>nexora_app</code> password from the migration default	Literal password in <code>0030_rls_hardening.sql</code>
Set <code>DATABASE_MIGRATION_URL</code> distinct from <code>DATABASE_URL</code>	Keeps runtime under RLS
Do not expose <code>/internal/metrics</code> or <code>/health/*</code> publicly	Unauthenticated by design
Bind the worker observability port to the cluster network only	ADR-026
Set an explicit <code>SERVER_CORS_ORIGINS</code> allowlist	Production rejects <code>*</code> but an empty list disables CORS entirely
Enable TLS termination upstream	The application serves plain HTTP
Review query-string API keys in proxy logs	Required by the CE contract on <code>/api/v2/ce/*</code>
Set <code>SERVER_TRUST_PROXY</code> appropriately	Affects client IP attribution in audit records

23 Troubleshooting

ChannelEngine and StockConnect issues have a dedicated table in that chapter. This section covers the rest of the system.

23.1 Authentication and authorization

Symptom	Cause	Investigation
401 on a valid token	Refresh session revoked, expired, or user has no effective permissions	<code>refresh_sessions</code> status; <code>AuthenticateAccessTokenUseCase</code> rejects empty permission sets
All refresh attempts fail after one succeeded	Token family reuse detected and the family revoked	Look for <code>REFRESH_REUSE_DETECTED</code> in the audit log; the client must re-authenticate
403 with a correct-looking role	Runtime checks are exact string matches; wildcards apply only at seeding	Compare the required key with <code>GET /api/v1/permissions</code>
API key works but lacks a new permission	Scopes were frozen at issuance	Rotate or reissue the key
Cannot remove the last admin role	Last-admin guard	Assign another admin first

23.2 Inventory and orders

Symptom	Cause	Investigation
Reservation fails with a business rule error	Insufficient <code>available</code>	<code>GET /api/v1/inventory/:productId</code> ; check active reservations
Order shows <code>SHIPPED</code> but nothing dispatched	Order status reflects full allocation, not dispatch	Read shipment status
Order never becomes <code>DELIVERED</code>	No code drives that transition	Use shipment status
Return receive rejected	The return was not approved first	Approve, then receive
Duplicate order from a marketplace	Expected; equivalence returns the existing order	Check for <code>ingestionOutcome: 'duplicate'</code>
Conflict on a repeated external reference	The payload differs from the stored order	Compare payloads; do not reuse a reference for different data

23.3 Synchronisation

Symptom	Cause	Investigation
Sync job never runs	Worker not running, or the event was not catalog-relevant	Worker logs, <code>queue_jobs_total</code>

Symptom	Cause	Investigation
Job skipped silently	Channel or marketplace inactive, offer has no external reference, or channel has no stock location	<code>ExecuteCatalogSyncJob</code> skip logs
Repeated retries then silence	Attempts exhausted	<code>QUEUE_DEFAULT_ATTEMPTS</code> ; inspect the adapter error
Permanent failure on first attempt	Adapter classified the error as auth, validation, conflict or unsupported	<code>mapMarketplaceErrorToAdapterError</code>
Product sync unsupported	Noon and Namshi throw by design	Create the SKU on the marketplace first
Inventory updates lag	Jobs coalesce by design	Expected; check the pending job count

23.4 Webhooks

Symptom	Cause	Investigation
Subscription rejected at creation	URL failed SSRF validation or event types not deliverable	HTTPS only, public address, allowlisted events
Deliveries stop after several attempts	Dead-lettered after exhausting attempts	<code>GET /api/v1/webhooks/:id/deliveries</code>
Signature verification fails at the receiver	Body must be verified exactly as transmitted	Header is <code>X-Nexora-Signature: v1=<hex></code> , HMAC-SHA256
Retries slower than expected	<code>Retry-After</code> is honoured up to the configured cap	<code>WEBHOOK_DELIVERY_MAX_RETRY_AFTER_SECONDS</code>
Inbound marketplace webhook returns 401	Ingress token unknown, or provider authentication failed	Token hash lookup, then the provider adapter

23.5 Platform

Symptom	Cause	Investigation
Startup fails at migration	Checksum mismatch on an applied migration	An applied file was edited; restore it and write a new migration
Startup hangs at migration	Another instance holds the advisory lock	Wait, or inspect <code>pg_locks</code> for <code>8147302915</code>
Queries fail with permission errors	Tenant GUC not set for an RLS-protected table	Ensure <code>database.execute(..., { tenantId })</code>
Readiness fails only for queue	Redis reachable but BullMQ unhealthy	Check <code>QUEUE_REDIS_URL</code> and the Redis eviction policy
Memory grows on the worker	Concurrency too high for job size	Lower <code>QUEUE_WORKER_CONCURRENCY</code>

24 Known limitations

Consolidated for planning. Each was verified in source; none is speculative.

24.1 Correctness and completeness

Area	Limitation
Order status	<code>DELIVERED</code> and <code>RETURNED</code> are defined but never set by any code path
Order status	<code>SHIPPED</code> indicates full shipment allocation, not physical dispatch
Cancellations	Always complete synchronously; <code>REJECTED</code> is unreachable through HTTP
Cancellations	<code>POST /api/v1/orders/:orderId/cancel</code> lacks the idempotency wrapper that <code>POST /api/v1/cancellations</code> has
Inventory	<code>TRANSFER_IN</code> and <code>TRANSFER_OUT</code> movement types have no use case
Inventory	Release reuses the <code>inventory.reserve</code> permission
Returns	Equivalence ignores per-line reason codes
Shipments	Equivalence compares carrier and tracking but not the service field
Storage	The S3 provider is wired and health-checked but unused by domain modules

24.2 Compatibility surface

Area	Limitation
CE catalog and delivery-state	Compute an idempotency key but never persist it
CE product push	Ensures a SKU exists; no rich content mapping
CE order poll	No <code>NEW</code> -only filter
CE invoice	Placeholder PDF with metadata only
CE delivery state	Only <code>DELIVERED</code> and <code>CLOSED</code> are mapped
CE channel identifiers	Require a numeric <code>external_reference</code>
CE webhook bridge	Selected by a magic <code>description</code> string
CE field coverage	Several accepted fields are ignored

24.3 Security

Area	Limitation
RLS consistency	Pre-tenant authentication paths query RLS tables without a tenant context
Secrets in migrations	<code>0030_rls_hardening.sql</code> contains a literal role password
MFA key default	<code>AUTH_MFA_ENCRYPTION_KEY</code> has a schema default
Login MFA	Not enforced at sign-in
Authorization	Exact matching only at runtime
API keys	Scopes frozen at issuance
CSP	Disabled in Helmet

24.4 Platform and process

Area	Limitation
CI	No workflow definitions in the repository
Liveness	API <code>/health/live</code> does not reflect shutdown
Configuration	<code>NOON_*</code> values are read but never populated
Rate limiting	The catalog sync limit is hard-coded at 120 per minute
Error codes	<code>PAYLOAD_TOO_LARGE</code> and <code>UNSUPPORTED_MEDIA_TYPE</code> are defined but never produced
Documentation drift	<code>tests/README.md</code> , the README phase table, and the Postgres port in the operations runbook are out of date

25 Future technical considerations

Recommendations only. Nothing here is implemented, and nothing here should be read as a commitment.

Theme	Suggestion	Rationale
RLS integrity	Give pre-tenant lookups an explicit, minimally privileged path rather than relying on RLS bypass	Removes the inconsistency between the restricted-role model and the code
CE idempotency parity	Route CE catalog and delivery-state commands through the idempotency service	Removes the only at-least-once side-effect path on the CE surface
CE catalog parity	Share one catalog command between <code>/api/v2</code> and <code>/api/v2/ce</code>	Eliminates a class of "works on one surface" defects
Order delivery status	Decide whether shipment delivery should advance the order aggregate	The state machine currently promises transitions nothing performs
CI automation	Add a pipeline running <code>verify</code> plus the integration suite against service containers	Makes the existing quality gates enforced rather than optional
API liveness	Make <code>/health/live</code> shutdown-aware, matching the worker	Removes an inconsistency between the two processes
Configurable sync limits	Move the catalog sync rate limit into configuration	Different marketplaces have different tolerances
Invoice documents	Implement the invoice port against real data	Replaces a placeholder with a usable document
Live CE verification	Complete the staging exercise the cutover runbook identifies	The only material gap in compatibility assurance
Documentation hygiene	Correct the drift listed under known limitations	Cheap, and removes misleading guidance

26 Glossary

Term	Meaning in Nexora
Adapter	Provider-specific implementation of a marketplace port
Available	<code>on_hand</code> - reserved ; the sellable quantity
Channel	A tenant's sales destination, referencing a marketplace
Channel-fulfilled order	An order the marketplace already shipped; no reservation is taken
ChannelEngine	Third-party channel manager whose API shape Nexora imitates; never called by Nexora
<code>compat_v2</code>	Provider namespace for compatibility integer identifiers
Compatibility surface	The <code>/api/v2</code> and <code>/api/v2/ce</code> ChannelEngine-shaped APIs
Effective price	The price applicable to a product on a channel at an instant
Equivalence	Comparing a create request against an existing aggregate to detect a safe duplicate
Idempotency key	Client-supplied header making a mutation safely retryable
Ingress token	Secret path segment authenticating an inbound marketplace webhook
Integration event	A domain occurrence written to the outbox for asynchronous delivery
Merchant SKU	Tenant-unique product code, the primary catalogue key
Minor units	Integer smallest currency unit; all money is stored this way
Movement	Immutable inventory ledger row
Offer	A product listed on a channel
Outbox	Table holding events committed with their domain change
Principal	The authenticated actor: a user or an API key
Readiness probe	A named dependency check contributing to <code>/health/ready</code>
Reservation	Stock held for an order but not yet consumed
RLS	PostgreSQL row-level security enforcing tenant isolation
Step-up	Short-lived elevated assurance obtained through MFA
StockConnect	External merchant system consuming the CE compatibility API
Tenant	The top-level isolation boundary

27 Appendix A — Feature to source map

Feature	Primary source
API process endpoint	<code>src/app/main.js</code>
Worker process endpoint	<code>src/workers/main.js</code>
Infrastructure composition	<code>src/app/bootstrap/create-infrastructure.js</code>
Application composition	<code>src/app/bootstrap/create-application.js</code>
Graceful shutdown	<code>src/app/bootstrap/shutdown.js</code>
Configuration schema	<code>src/app/config/schema.js</code> , <code>src/app/config/config.js</code>
HTTP server and plugin order	<code>src/app/http/create-server.js</code>
Authentication	<code>src/app/http/plugins/authentication.plugin.js</code>
Error normalisation and envelope	<code>src/app/errors/error-mapper.js</code> , <code>src/app/errors/error-envelope.js</code>
Error hierarchy	<code>src/shared/errors/app-error.js</code> , <code>error-codes.js</code>
Readiness	<code>src/app/observability/readiness.js</code>
Request context	<code>src/shared/context/request-context.js</code>
Idempotency fingerprint	<code>src/shared/idempotency/fingerprint.js</code>
Event catalogue	<code>src/shared/events/event-catalog.js</code>
Money value object	<code>src/shared/money/money.js</code>
Rate limit policies	<code>src/shared/auth/rate-limit-policies.js</code>
CE prefix detection	<code>src/shared/auth/merchant-compat-route-prefix.js</code>
CE credential readers	<code>src/shared/auth/read-merchant-compat-query-api-key.js</code>
Merchant compatibility routes	<code>src/modules/compatibility/presentation/compatibility.routes.js</code>
StockConnect CE routes	<code>src/modules/compatibility/presentation/stockconnect-ce.routes.js</code>
CE error envelope	<code>src/modules/compatibility/application/errors/map-core-error.js</code>
CE idempotency derivation	<code>src/modules/compatibility/application/resolve-stockconnect-ce-idempotency-key.js</code>
CE catalog command	<code>src/modules/compatibility/application/stockconnect-ce-catalog-command.js</code>
CE order poll	<code>src/modules/compatibility/application/stockconnect-ce-order-compatibility-query.js</code>
CE invoice generation	<code>src/modules/compatibility/application/generate-stockconnect-ce-invoice-pdf.js</code>

Feature	Primary source
CE readiness probe	<code>src/modules/compatibility/application/stockconnect-ce-readiness.js</code>
CE webhook body	<code>src/modules/compatibility/application/build-stockconnect-ce-webhook-body.js</code>
CE webhook constants	<code>src/shared/stockconnect/stockconnect-ce-webhook-subscription.js</code>
External identifier resolution	<code>src/modules/compatibility/application/compatibility-external-id-resolution.js</code>
Order status machine	<code>src/modules/orders/domain/order-status.js</code>
Order fulfilment coordination	<code>src/modules/orders/application/order-fulfillment-service.js</code>
Channel order equivalence	<code>src/modules/orders/application/channel-order-request-equivalence.js</code>
Inventory balance invariant	<code>src/modules/inventory/domain/inventory-balance.js</code>
Shipment inventory consumption	<code>src/modules/shipments/application/fulfill-inventory-when-shipment-shipped.js</code>
Catalog sync planning	<code>src/modules/channel-catalog-sync/application/plan-catalog-sync-jobs.js</code>
Catalog sync execution	<code>src/modules/channel-catalog-sync/application/execute-catalog-sync-job.js</code>
Catalog sync job identifiers	<code>src/modules/channel-catalog-sync/application/build-catalog-sync-job-id.js</code>
Adapter runtime factory	<code>src/modules/marketplaces/application/marketplace-adapter-runtime-factory.js</code>
Adapter error classification	<code>src/modules/marketplaces/application/map-marketplace-error-to-adapter-error.js</code>
Inbound webhook receipt	<code>src/modules/marketplace-webhook-ingestion/application/receive-marketplace-webhook.js</code>
Order ingestion service	<code>src/modules/marketplace-order-ingestion/application/marketplace-order-ingestion-service.js</code>
Webhook delivery	<code>src/modules/webhooks/application/webhook-delivery-service.js</code>
Webhook signing	<code>src/modules/webhooks/application/webhook-request-signer.js</code>
Webhook response classification	<code>src/modules/webhooks/application/classify-webhook-http-response.js</code>
SSRF validation	<code>src/shared/security/validate-outbound-https-url.js</code> , <code>ssrf-validator.js</code>
Postgres pool and transactions	<code>src/infrastructure/postgres/postgres-database.js</code>
Migration runner	<code>src/infrastructure/postgres/migrator.js</code>
Queue names and jobs	<code>src/infrastructure/queue/queue-names.js</code>
Worker runtime	<code>src/infrastructure/queue/bullmq-worker-runtime.js</code>
Worker job handlers	<code>src/workers/handlers/queue-job-handlers.js</code>

Feature	Primary source
Rate limiter	<code>src/infrastructure/redis/redis-rate-limiter.js</code>
Outbound HTTP client	<code>src/infrastructure/http/fetch-http-client.js</code>
Logger and redaction	<code>src/infrastructure/observability/pino-logger.js</code> , <code>src/shared/logging/redaction.js</code>
Metrics	<code>src/infrastructure/observability/prometheus-metrics.js</code>
Tracing	<code>src/infrastructure/observability/tracing.js</code>
Architecture rules	<code>.dependency-cruiser.cjs</code>

28 Appendix B — Architecture decision records

ADR	Title	Status
001	Modular monolith	Accepted
002	PostgreSQL as source of truth	Accepted
003	Redis infrastructure layer	Accepted
004	Object storage	Accepted
005	Transactional outbox pattern	Accepted
006	Distributed rate limiting	Accepted
007	ChannelEngine compatibility layer	Accepted, scope in ADR-018
008	Multi-tenancy	Accepted
009	Audit log architecture	Accepted
010	API versioning strategy	Accepted
011	Database access with node-postgres	Accepted
012	Standard system roles	Accepted
013	Inventory concurrency and authority	Accepted
014	Pricing validity and resolution	Accepted
015	Offer lifecycle	Accepted
016	Phase 4 commerce public contracts	Accepted
017	Phase 4 orders and fulfilment	Accepted
018	Phase 5 merchant-compatible API scope	Accepted
019	Phase 6 integration events and external webhooks	Accepted
020	Phase 7 channel and marketplace inbound integration	Accepted
021	External integer ID mapping	Accepted
022	Phase 9 channel default stock location	Accepted
023	Phase 10 webhook delivery history retention	Accepted
024	Phase 11 commerce webhook catalogue expansion	Accepted
025	Phase 12 webhook Retry-After scheduling	Accepted
026	Phase 13 worker observability HTTP surface	Accepted

ADR	Title	Status
027	Phase 14 shipment fulfilment inventory integration	Accepted, specification
028	Phase 15 product, inventory and pricing sync architecture	Accepted, specification
029	Marketplace connector framework	Accepted
030	Marketplace connection and entity mapping	Accepted
031	Phase 32 generic marketplace webhook framework	Accepted

ADR-018 predates the current code in one respect: it describes the initial `/api/v2` surface as Merchant-only without `POST /v2/orders`, whereas channel order ingest is now implemented there. The ADR record has not been superseded.

29 Appendix C — Related documentation

Document	Contents
<code>docs/API_ROUTE_INVENTORY.md</code>	Generated inventory of all 147 endpoints with route identifiers
<code>docs/SWAGGER_COVERAGE_REPORT.md</code>	Route-to-OpenAPI coverage verification
<code>docs/POSTMAN_API_INVENTORY.md</code>	Inventory backing the Postman collection
<code>docs/POSTMAN_COVERAGE_REPORT.md</code>	Route-to-Postman request mapping
<code>docs/POSTMAN_USAGE.md</code>	How to import and use the collection
<code>docs/architecture/compatibility-matrix.md</code>	CE endpoint parity matrix
<code>docs/architecture/compatibility-idempotency.md</code>	Idempotency ledger semantics
<code>docs/architecture/stockconnect-ce-operations.md</code>	CE operational guidance
<code>docs/architecture/stockconnect-ce-production-cutover.md</code>	Cutover runbook and verification status
<code>docs/architecture/platform-independence.md</code>	Provider-neutrality rationale
<code>docs/marketplaces/capability-matrix.md</code>	Per-marketplace capability matrix
<code>docs/marketplaces/order-ingestion.md</code>	Order ingestion detail
<code>docs/marketplaces/webhook-ingestion.md</code>	Inbound webhook detail
<code>docs/operations/README.md</code>	Operational runbook
<code>docs/decisions/</code>	The full ADR set